

DSC 102

Systems for Scalable Analytics

Fall 2023

Haojian Jin

Now for the course logistics ...

Prerequisites

- ❖ **DSC 100** (or equivalent) is necessary
- ❖ Transitively **DSC 80**; a mainstream ML algorithmics course is necessary
- ❖ Proficiency in Python programming
- ❖ For all other cases, email me with proper justification; a waiver can be considered

<https://haojian.github.io/DSC102FA23/>

Components and Grading

- ❖ **3 Programming Assignments: 40% (8% + 16% + 16%)**
 - ❖ No late days! Plan your work well ahead.
 - ❖ **Plan your credit as well!**
- ❖ **Midterm Exam: 15%**
 - ❖ TBD; in-class only (50min)
- ❖ **Cumulative Final Exam: 35%**
 - ❖ Dec. 15; Canvas Quiz only (3hrs long but 4hrs limit)
- ❖ **10 (of 12) Peer Instruction Activities: 10%**
- ❖ **Extra Credit Peer Evaluation Activities: 4% (likely)**
- ❖ LMK ahead of time if you need makeup exam slot

<https://haojian.github.io/DSC102FA23/>

Grading Scheme

Hybrid of relative and absolute; grade is better of the two

Grade	Relative Bin (Use strictest)	Absolute Cutoff (>=)
A+	Highest 5%	95
A	Next 10% (5-15)	90
A-	Next 15% (15-30)	85
B+	Next 15% (30-45)	80
B	Next 15% (45-60)	75
B-	Next 15% (60-75)	70
C+	Next 5% (75-80)	65
C	Next 5% (80-85)	60
C-	Next 5% (85-90)	55
D	Next 5% (90-95)	50

Example: Score 82 but 33%ile; Rel.: B-; Abs.: B+; so, B+

Programming Assignments

- ❖ **PA0: Setting up AWS and Dask**
 - ❖ Oct 6 to Oct 22
- ❖ **PA1: Data Exploration with Dask**
 - ❖ Oct 23 to Nov 12
- ❖ **PA2: Feature Eng. and Model Selection with Spark**
 - ❖ Nov 13 to Dec 04
- ❖ **Expectations on the PAs:**
 - ❖ Teams of 1-3; see webpage on academic integrity
 - ❖ I will cover the concepts and tools' tradeoffs in the lectures
 - ❖ TAs will explain and demo the tools; handle all Q&A
 - ❖ You are expected to put in the effort to learn the details of the tools' APIs using their documentation on your own!

<https://haojian.github.io/DSC102FA23/>

Course Administrivia

- ❖ **Lectures:** MWF 3pm-3:50pm PT at **CENTR 214**
 - ❖ Attendance optional but encouraged; podcast available
 - ❖ No need for clickers.
- ❖ **Discussions:** Friday 9-9:50am PT on Zoom
 - ❖ Only for talks on PAs by TAs, for pre-exam review by me
- ❖ **Instructor:** Haojian Jin; haojian@ucsd.edu
 - ❖ OHs: **Wednesday 4-5 pm PT at 341 HDSI**
- ❖ **Slack** for all communications
- ❖ **Canvas** for PA submission, Peer Evaluation Activities, Final Exam

<https://haojian.github.io/DSC102FA23/>

Office hours

- ❖ Rohit Ramaprasad's OHs: Tuesday 1:00 PM - 2:00 PM
- ❖ Sai Sree Harsha's OHs: Thursday 1:00 PM - 2:00 PM
- ❖ Golokesh Patra's OHs: Monday 11:00 AM - 12:00 PM
- ❖ Post questions to the ta-public channel.
- ❖ Avoid asking repetitive questions.

<https://haojian.github.io/DSC102FA23/>

General Dos and Do NOTs

Do:

- ❖ Follow all announcements on Piazza
- ❖ Try to join the lectures/discussions live
- ❖ Raise your hand before speaking
- ❖ View/review podcast videos asynchronously by yourself
- ❖ To contact me/TAs, use private Slack; if you really need to email, use “DSC 102:” as subject prefix

Do NOT:

- ❖ Harass, intimidate, or intentionally talk over others
- ❖ Violate academic integrity on the PAs, exams, or other components; I am *very strict* on this matter!

Reasonable person.

- (1) Everyone will be reasonable.
- (2) Everyone expects everyone else to be reasonable.
- (3) No one is special.
- (4) Do not be offended if someone suggests you are not being reasonable.

Now for the course structure ...

DSC 102 will get you thinking about the **fundamentals of systems for scalable analytics**

1. “**Systems**”: What resources does a computer have? How to store and efficiently compute over large data? What is cloud?
2. “**Scalability**”: How to scale and parallelize data-intensive computations?
3. **For “Analytics”:**
 1. **Source**: Data acquisition & preparation for ML
 2. **Build**: Model selection & deep learning systems
 3. **Deploying** ML models
4. Hands-on experience with scalable analytics tools

Data Systems Concerns in ML

Key concerns in ML:

Q: How do “ML Systems” relate to ML?

Runtime efficiency (sometimes)

Additional key *practical* concerns in ML Systems:
ML Systems : ML :: Computer Systems : TCS

Scalability (and **efficiency** at scale)

Usability

Manageability

Developability

*Long-standing
concerns in the
DB systems
world!*

Q: Q: What if I didn't have the discipline to take my ideas from the PPT to code?

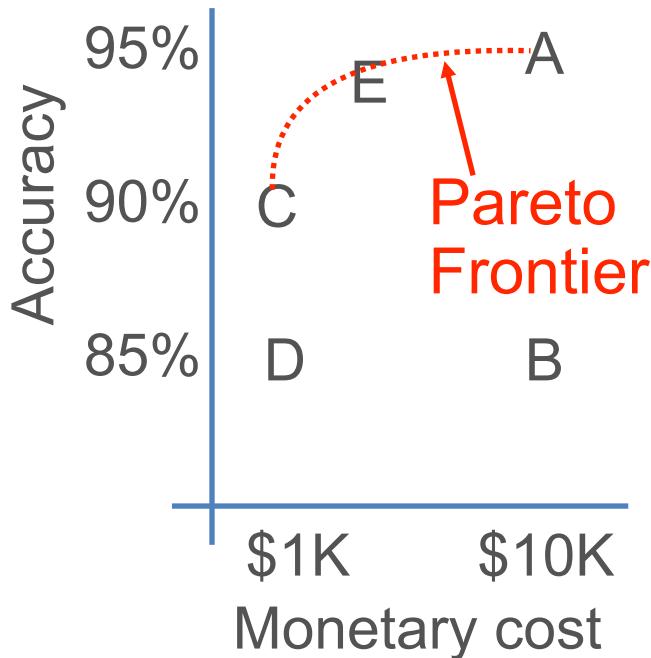
Conceptual System Stack Analogy

	Relational DB Systems	ML Systems
Theory	First-Order Logic Complexity Theory	Learning Theory Optimization Theory
Program Formalism	Relational Algebra	Tensor Algebra Gradient Descent
Program Specification	SQL	TensorFlow? Scikit-learn?
Program Modification	Query Optimization	???
Execution Primitives	Parallel Relational Operator Dataflows	Depends on ML Algorithm
Hardware	CPU, GPU, FPGA, NVM, RDMA, etc.	

Real-World ML: Pareto Surfaces

Q: Suppose you are given ad click-through prediction models A, B, C, and D with accuracies of 95%, 85%, 90%, and 85%, respectively. Which one will you pick?

Q: What about now?



- ❖ Real-world ML users must grapple with multi-dimensional *Pareto surfaces*: accuracy, monetary cost, training time, scalability, inference latency, tool availability, interpretability, fairness, etc.
- ❖ *Multi-objective optimization* criteria set by application needs / business policies.

Learning Outcomes of this course

- ❖ **Explain** the basic principles of the memory hierarchy, parallelism paradigms, scalable data systems, and cloud computing.
- ❖ **Identify** the abstract data access patterns of, and opportunities for parallelism and efficiency gains in, data processing and ML algorithms at scale.
- ❖ **Outline** how to use cluster and cloud services, dataflow (“Big Data”) programming with MapReduce and Spark, and ML tools at scale.
- ❖ **Apply** the above programming skills to create end-to-end pipelines for data preparation, feature engineering, and model selection on large-scale datasets.
- ❖ **Reason** critically about practical tradeoffs between accuracy, runtimes, scalability, usability, and total cost.

What this course is NOT about

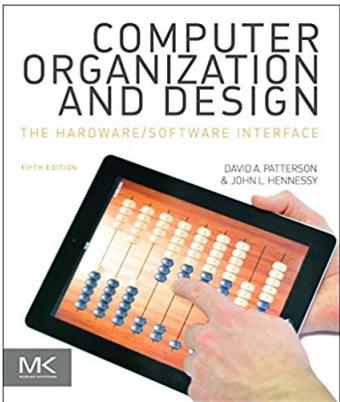
- ❖ NOT a course on databases, relational model, or SQL
 - ❖ Take DSC 100 instead (pre-requisite)
- ❖ NOT a course on internal details of RDBMSs
 - ❖ Take CSE 132C instead
- ❖ NOT a training module for how to use Spark
- ❖ NOT a course on ML or data mining *algorithmics*; instead, we focus on ML *systems*

Tentative Course Schedule

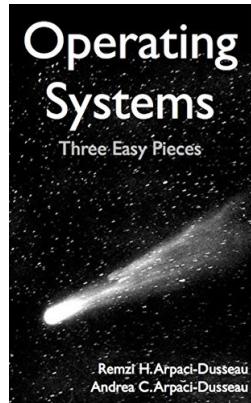
Week	Topic
Systems Principles 4	Basics of Machine Resources: Computer Organization Basics of Machine Resources: Operating Systems Basics of Cloud Computing
Scalability Principles 4-5	Parallel and Scalable Data Processing: Parallelism Basics Midterm Exam on TBD
6-7	Parallel and Scalable Data Processing: Scalable Data Access
7-8	Parallel and Scalable Data Processing: Data Parallelism
9	Dataflow Systems
10	ML Model Building Systems
11	Final Exam on Dec 15

There will be 2 industry guest lectures (maybe 3)

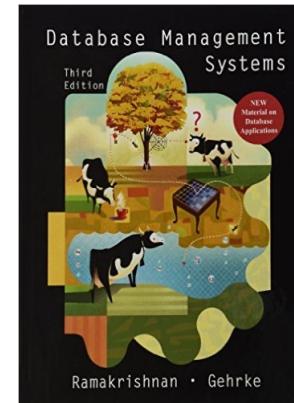
Suggested Textbooks



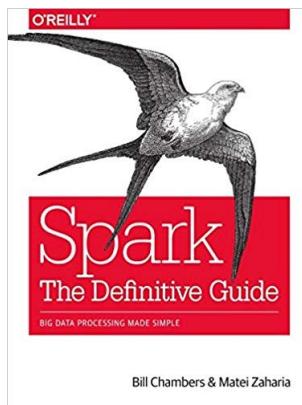
Aka “CompOrg Book”



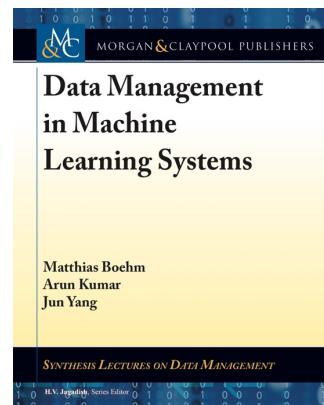
Aka “Comet Book”



Aka “Cow Book”



Aka “Spark Book”

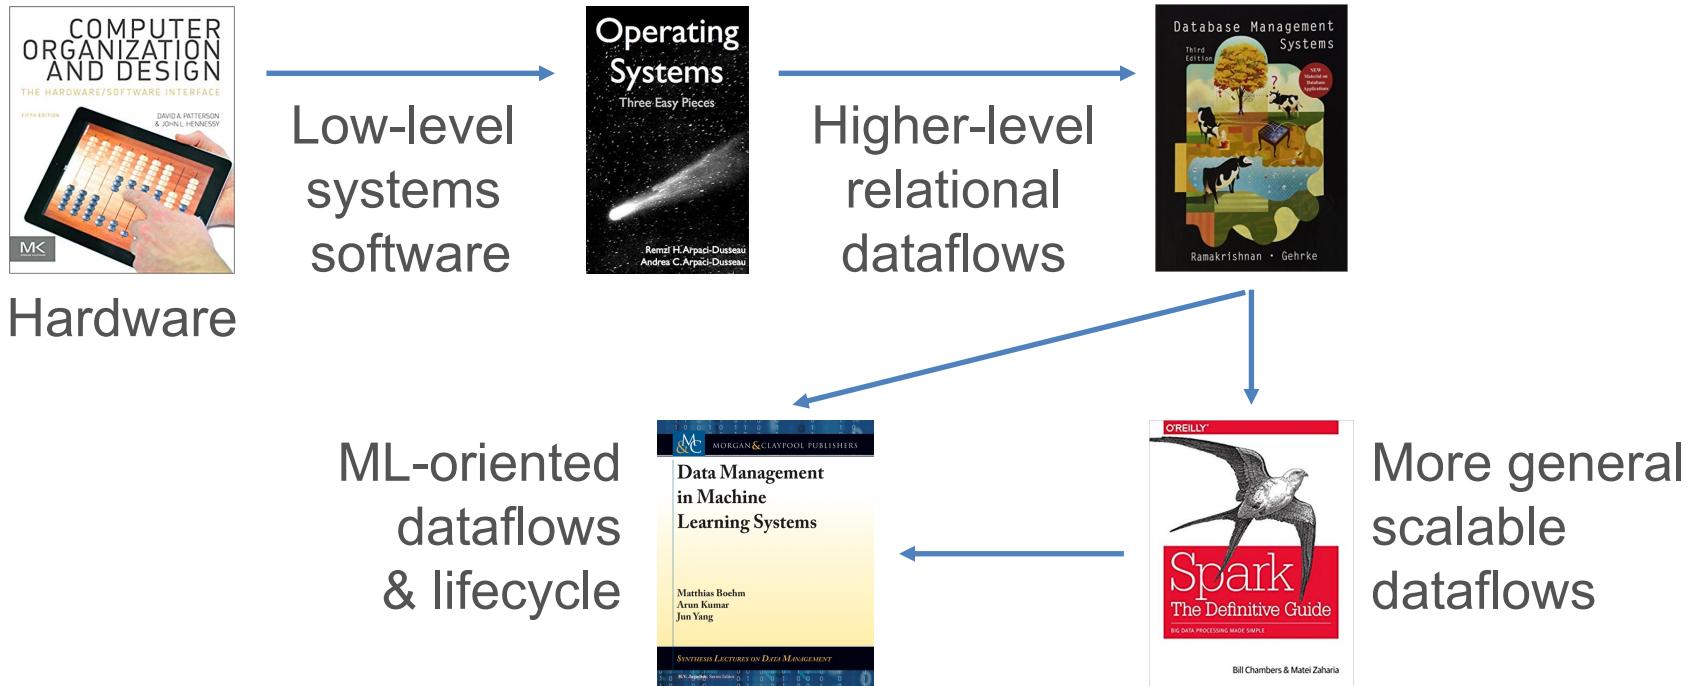


Aka “MLSys Book”

(Free PDFs available online; also check out our library)

Why so many textbooks?!

1. Computer systems are about carefully layering *levels of abstraction*.



2. Analytics/ML Systems is a recent/emerging area of research.
3. Also, DSC 102 is the first UG course of its kind in the world!

Tentative Course Schedule

Week	Topic
1-2	Basics of Machine Resources: Computer Organization
Systems Principles	Basics of Machine Resources: Operating Systems
	Basics of Cloud Computing
4-5	Parallel and Scalable Data Processing: Parallelism Basics
6	Midterm Exam on TBD
6-7	Parallel and Scalable Data Processing: Scalable Data Access
7-8	Parallel and Scalable Data Processing: Data Parallelism
9	Dataflow Systems
10	ML Model Building Systems
11	Final Exam on Dec 15

There will be 2 industry guest lectures (maybe 3)

DSC 102

Systems for Scalable Analytics

Topic 1: Basics of Machine Resources
Part 1: Computer Organization

Ch. 1, 2.1-2.3, 2.12, 4.1, and 5.1-5.5 of CompOrg Book

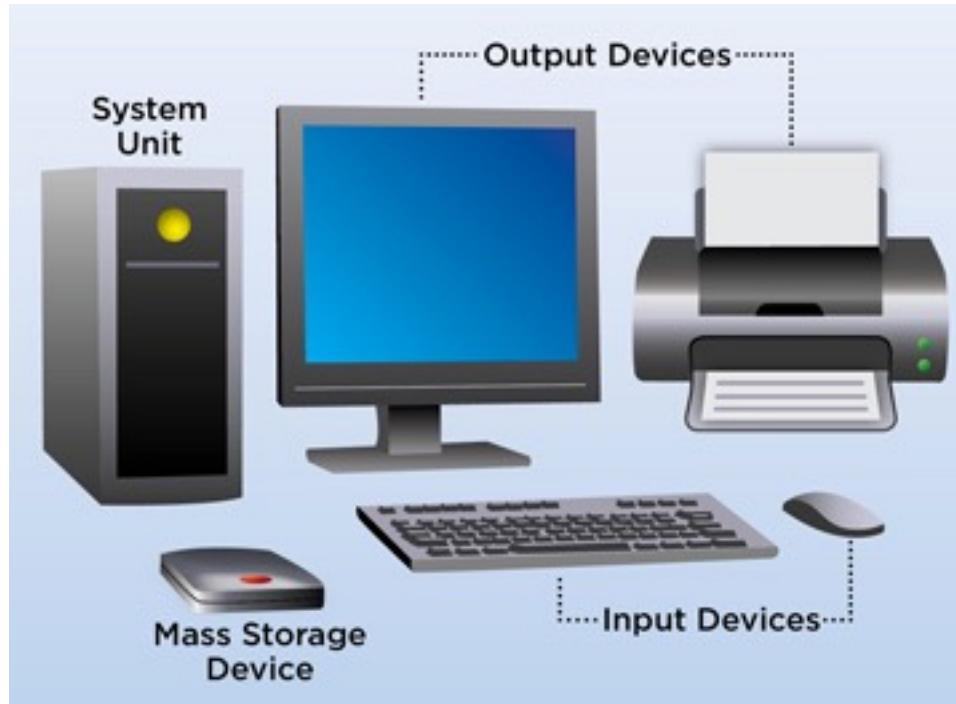
Q: What is a computer?

A programmable electronic device that can store, retrieve,
and process digital data.

Outline

- ➔ ♦ Basics of Computer Organization
 - ♦ Digital Representation of Data
 - ♦ Processors and Memory Hierarchy
- ♦ Basics of Operating Systems
 - ♦ Process Management: Virtualization; Concurrency
 - ♦ Filesystem and Data Files
 - ♦ Main Memory Management
- ♦ Persistent Data Storage

Parts of a Computer



Hardware:

The electronic machinery (wires, circuits, transistors, capacitors, devices, etc.)

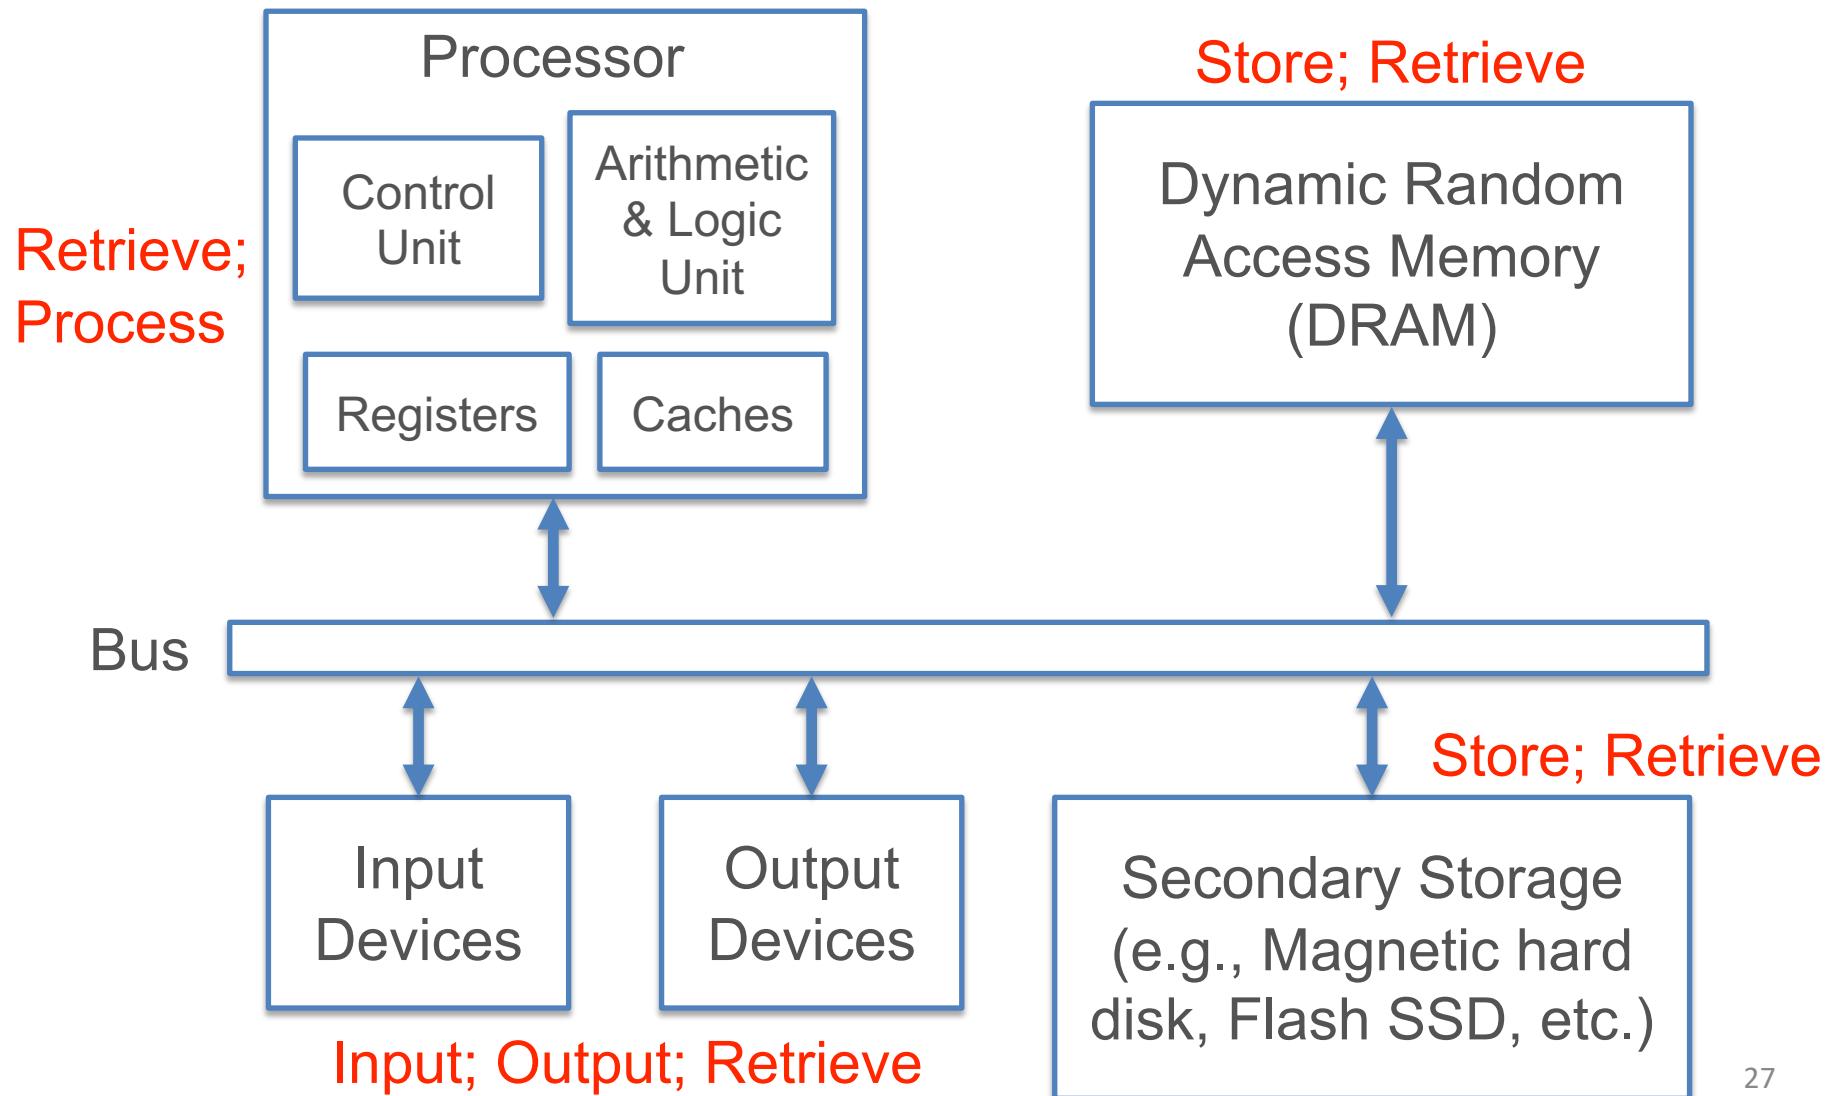
Software:

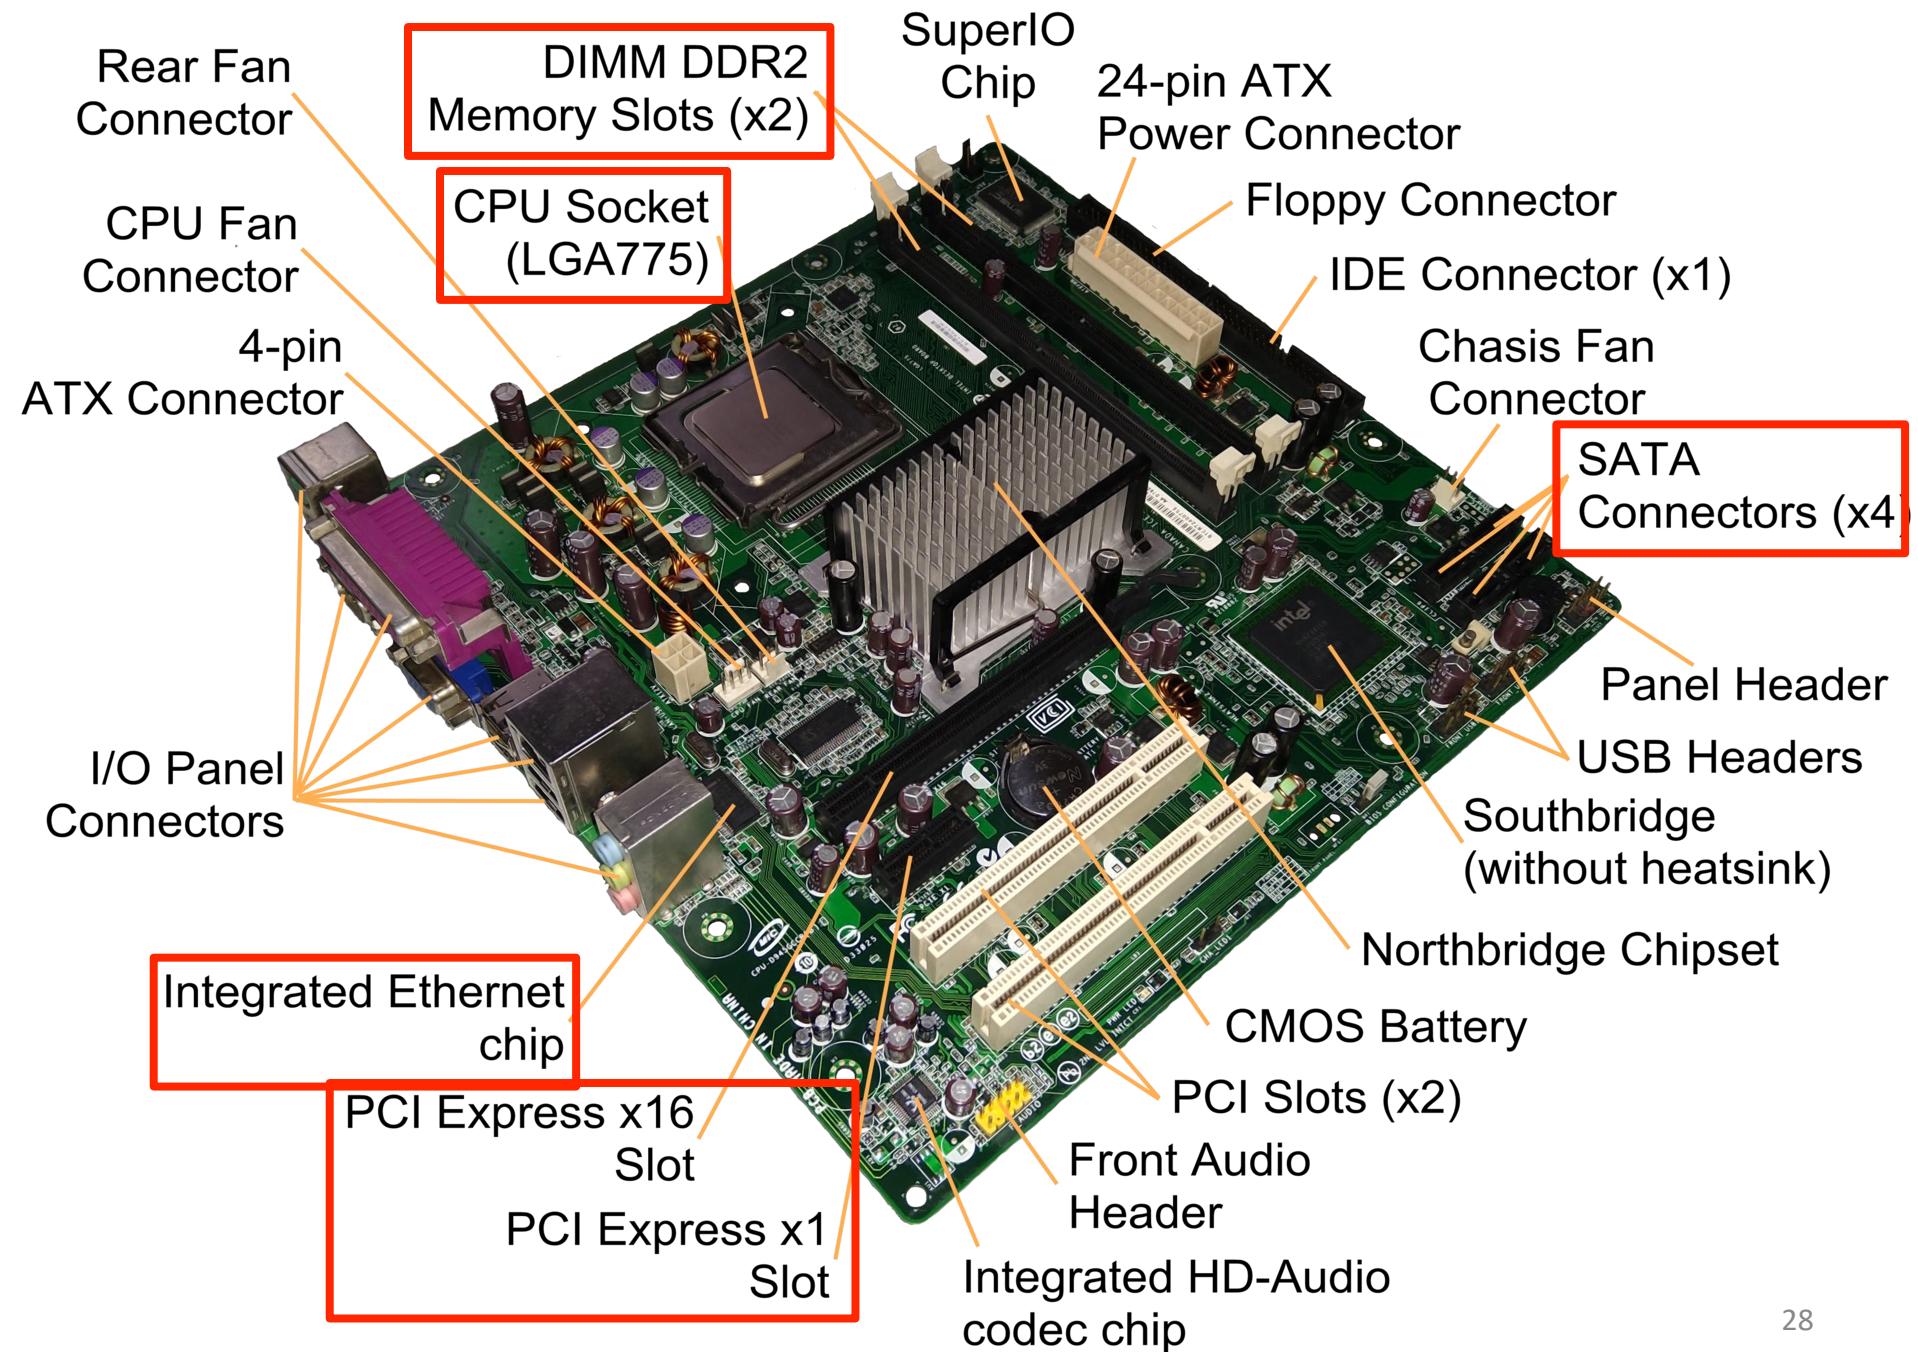
Programs (instructions) and data

Key Parts of Computer Hardware

- ❖ **Processor** (CPU, GPU, etc.)
 - ❖ Hardware to orchestrate and execute *instructions* to manipulate *data* as specified by a *program*
- ❖ **Main Memory** (aka Dynamic Random Access Memory)
 - ❖ Hardware to store *data* and *programs* that allows very fast location/retrieval; byte-level addressing scheme
- ❖ **Disk** (aka secondary/persistent storage)
 - ❖ Similar to memory but *persistent*, *slower*, and higher capacity / cost ratio; various addressing schemes
- ❖ **Network interface controller (NIC)**
 - ❖ Hardware to send data to / retrieve data over network of interconnected computers/devices

Abstract Computer Parts and Data





Key Aspects of Software

- ❖ **Instruction**
 - ❖ A command understood by hardware; finite vocabulary for a processor: Instruction Set Architecture (ISA); bridge between hardware and software
- ❖ **Program (aka code)**
 - ❖ A collection of instructions for hardware to execute
- ❖ **Programming Language (PL)**
 - ❖ A human-readable *formal* language to write programs; at a much higher level of *abstraction* than ISA
- ❖ **Application Programming Interface (API)**
 - ❖ A set of functions (“interface”) exposed by a program/set of programs for use by humans/other programs
- ❖ **Data**
 - ❖ Digital representation of *information* that is stored, processed, displayed, retrieved, or sent by a program

Main Kinds of Software

- ❖ **Firmware**
 - ❖ Read-only programs “baked into” a device to offer basic hardware control functionalities
- ❖ **Operating System (OS)**
 - ❖ Collection of interrelated programs that work as an intermediary platform/service to enable application software to use hardware more effectively/easily
 - ❖ Examples: Linux, Windows, MacOS, etc.
- ❖ **Application Software**
 - ❖ A program or a collection of interrelated programs to manipulate data, typically designed for human use
 - ❖ Examples: Excel, Chrome, PostgreSQL, etc.

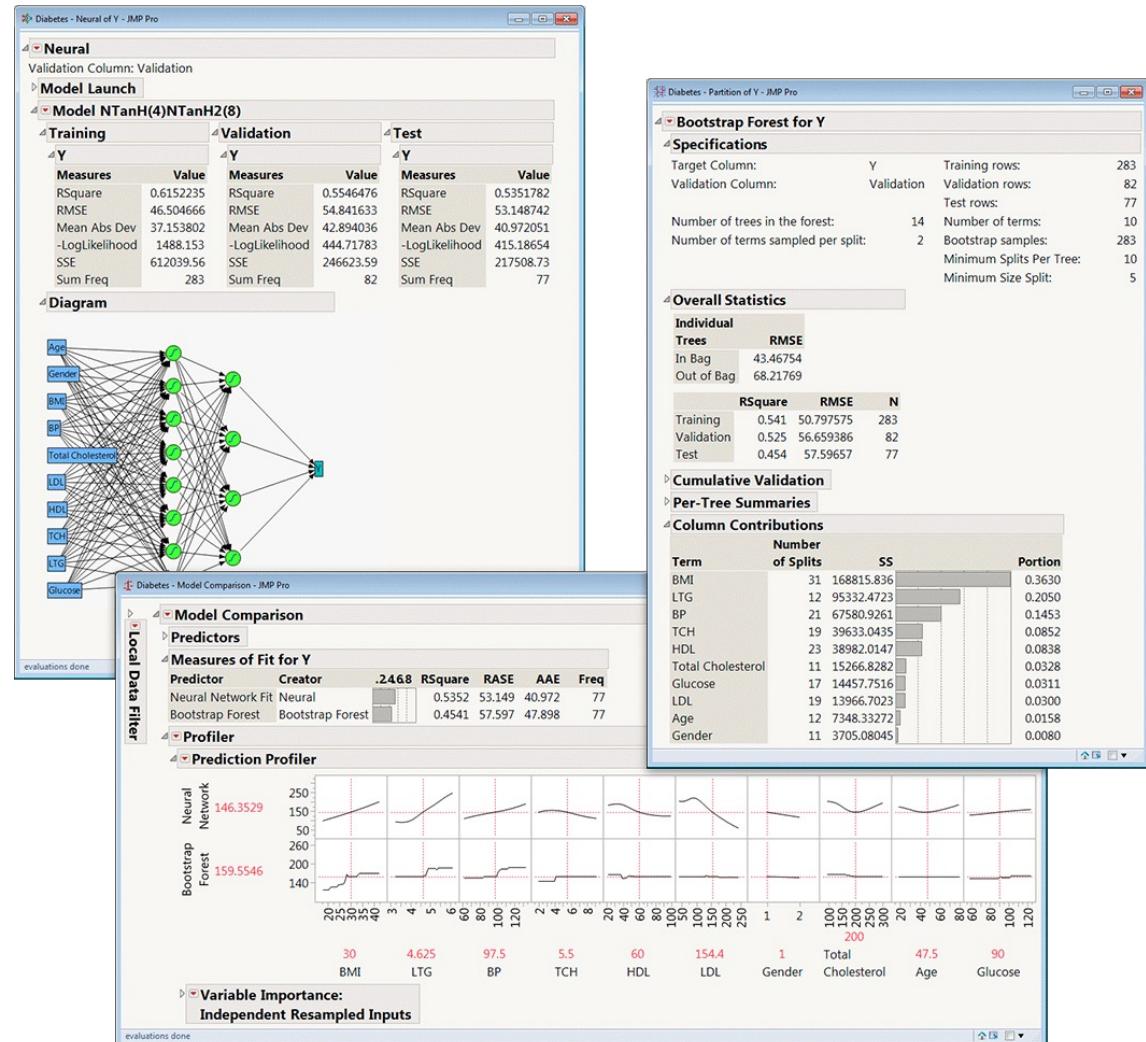
Outline

- ❖ Basics of Computer Organization
 - ❖ Digital Representation of Data
 - ❖ Processors and Memory Hierarchy
- ❖ Basics of Operating Systems
 - ❖ Process Management: Virtualization; Concurrency
 - ❖ Filesystem and Data Files
 - ❖ Main Memory Management
- ❖ Persistent Data Storage

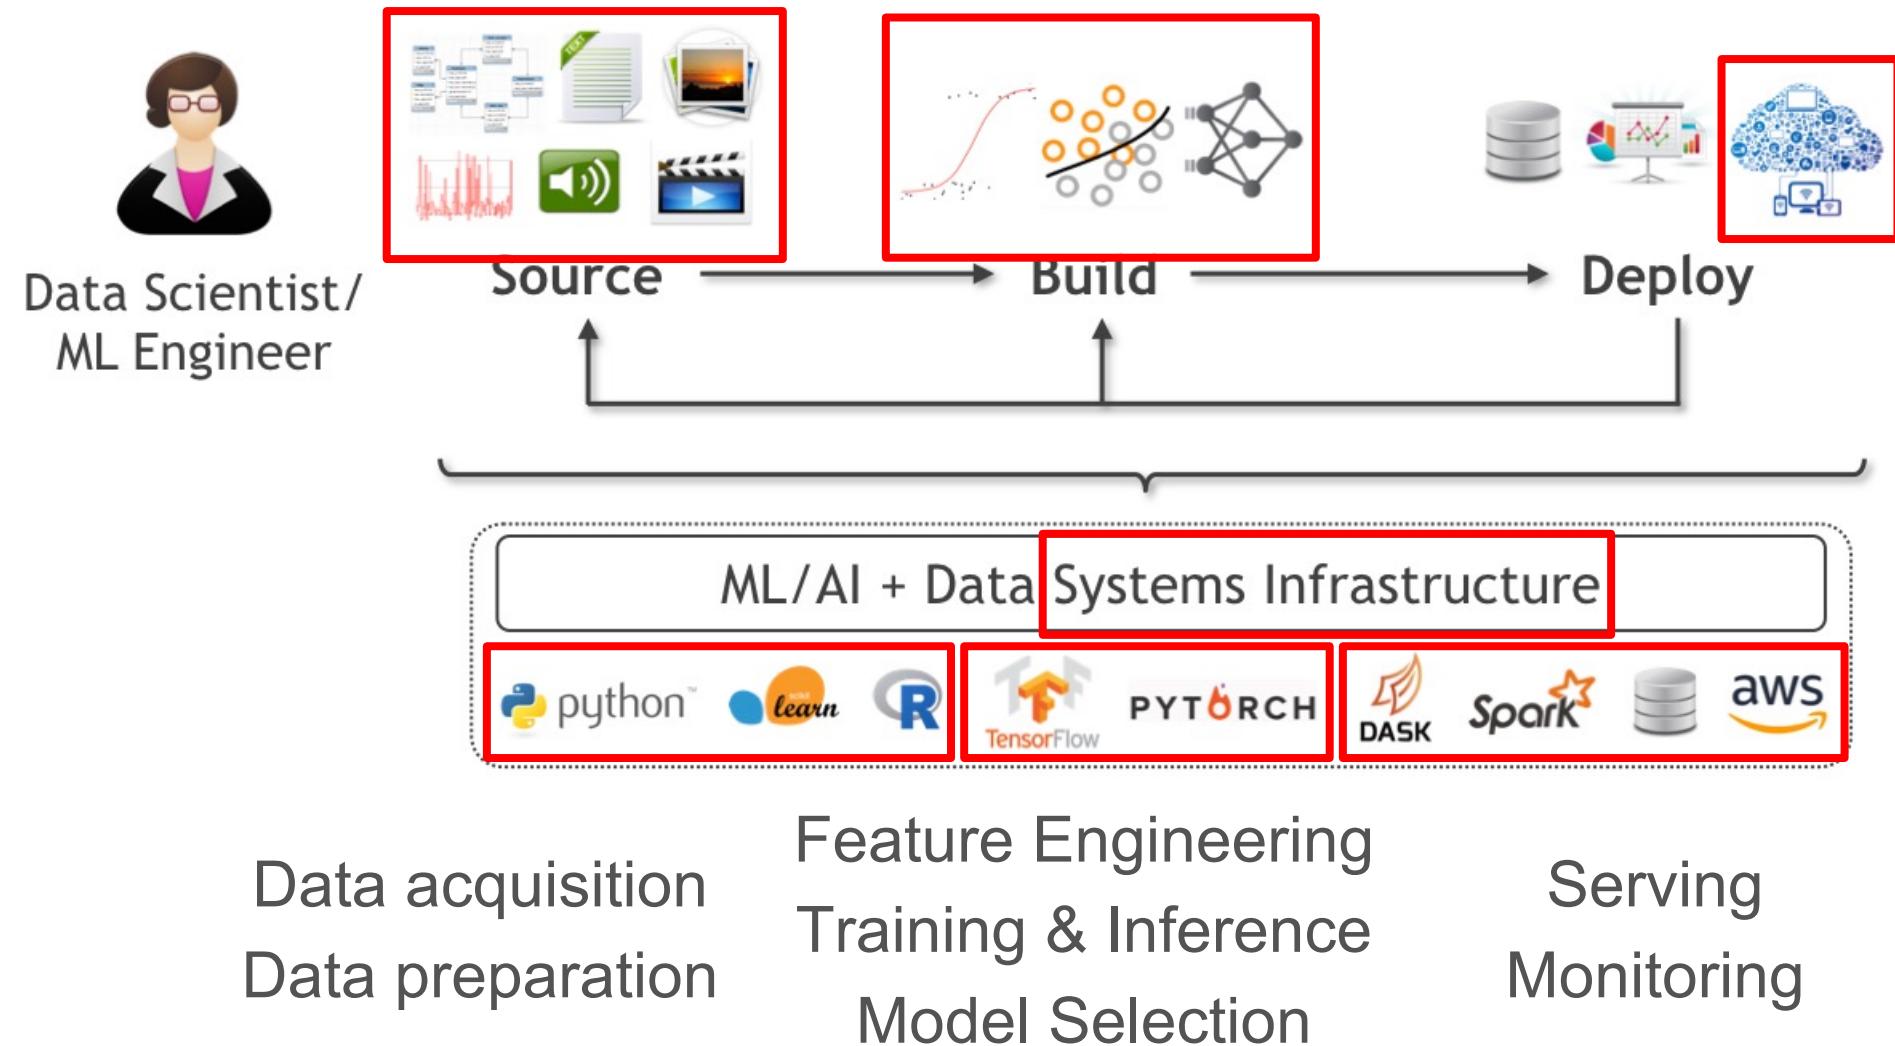
Q: But why bother learning such low-level computer sciencey stuff in Data Science?

Luxury of “Statisticians”/“Analysts” of Yore

- ❖ **Methods:** Sufficed to learn just math/stats, maybe some SQL
- ❖ **Types:** Mostly tabular (relational), maybe some time series
- ❖ **Scale:** Mostly small (KBs to few GBs)
- ❖ **Tools:** Simple GUIs for both analysis and deployment; maybe an R-like console



Reality of Today's “Data Scientists”



Why bother with these in Data Science?

- ❖ Basics of Computer Organization
 - ❖ Digital Representation of Data
 - ❖ Processors and Memory Hierarchy
 - ❖ Basics of Operating Systems
 - ❖ Process Management: Virtualization; Concurrency
 - ❖ Filesystem and Data Files
 - ❖ Main Memory Management
 - ❖ Persistent Data Storage
- You will face myriad and new data types
- Compute hardware is evolving fast
- You will need to use new methods on evolving data file formats on clusters / cloud
- Storage hardware is evolving fast



statistician

Location



Statistician Salaries United States ▾

Overview

Salaries

Interviews

Insights

Career Path

How much does a Statistician make?

Updated Jan 4, 2022

Industry

Employer Size

Experience

All industries

All company sizes

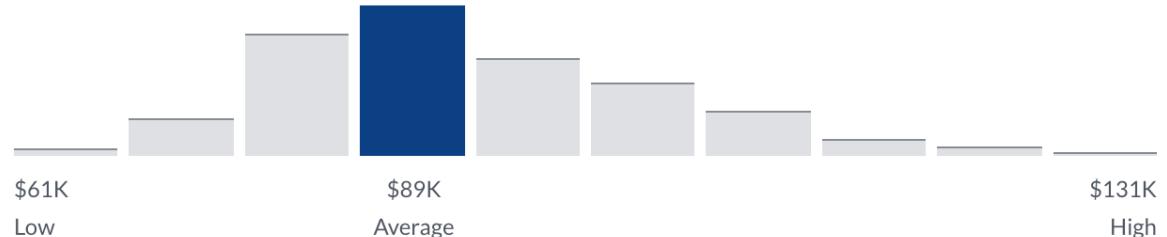
All years of Experience

Very High Confidence

\$88,989 /yr

Average Base Pay

2,398 salaries





Data Scientist Salaries United States ▾

Overview

Salaries

Interviews

Insights

Career Path

How much does a Data Scientist make?

Updated Jan 4, 2022

Industry

▼

Employer Size

▼

Experience

▼

To filter salaries for Data Scientist, [Sign In](#) or [Register](#).

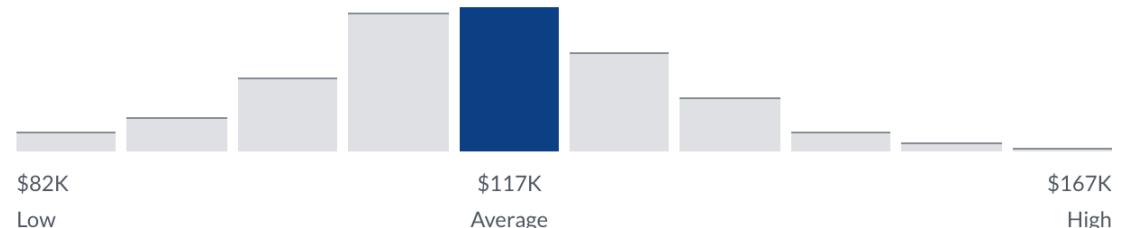


Very High Confidence

\$117,212 /yr

Average Base Pay

18,354 salaries

**— 88,989****= 28,223!**

Outline

- ❖ Basics of Computer Organization
- ➔ ❖ Digital Representation of Data
 - ❖ Processors and Memory Hierarchy
- ❖ Basics of Operating Systems
 - ❖ Process Management: Virtualization; Concurrency
 - ❖ Filesystem and Data Files
 - ❖ Main Memory Management
- ❖ Persistent Data Storage

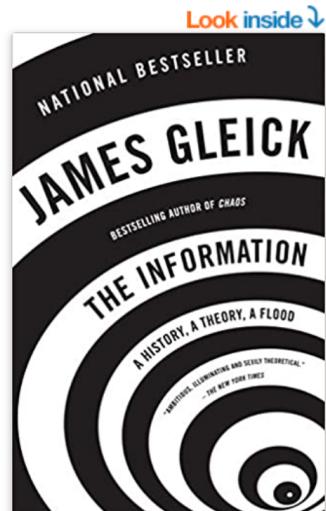
Q: What is data?

上 5 もう 14 年も前から、この工場で作られる車の品質は、常に高水準で、多くの車好きたちに愛されています。しかし、最近では、車の品質が下がり始めています。そこで、車の品質を回復するための取り組みが行われています。

車の品質を回復するための取り組みには、以下の点があります。

1. 車の品質を回復するための取り組みには、以下の点があります。

Why bits?



Listen



[See this image](#)

The Information: A History, A Theory, A Flood Paperback – Illustrated, March 6, 2012

by James Gleick (Author)

4.5 out of 5 stars 1,339 ratings

Editors' pick Best Science Fiction & Fantasy

Searching ••• S +

[See all formats and editions](#)

Kindle

\$9.99

[Read with Our Free App](#)

Audiobook

\$0.00

[Free with your Audible trial](#)

Hardcover

\$13.25 - \$39.94

Other new, used and
collectible from \$2.39

Paperback

\$12.23

Other new and used
from \$2.45

Audio CD

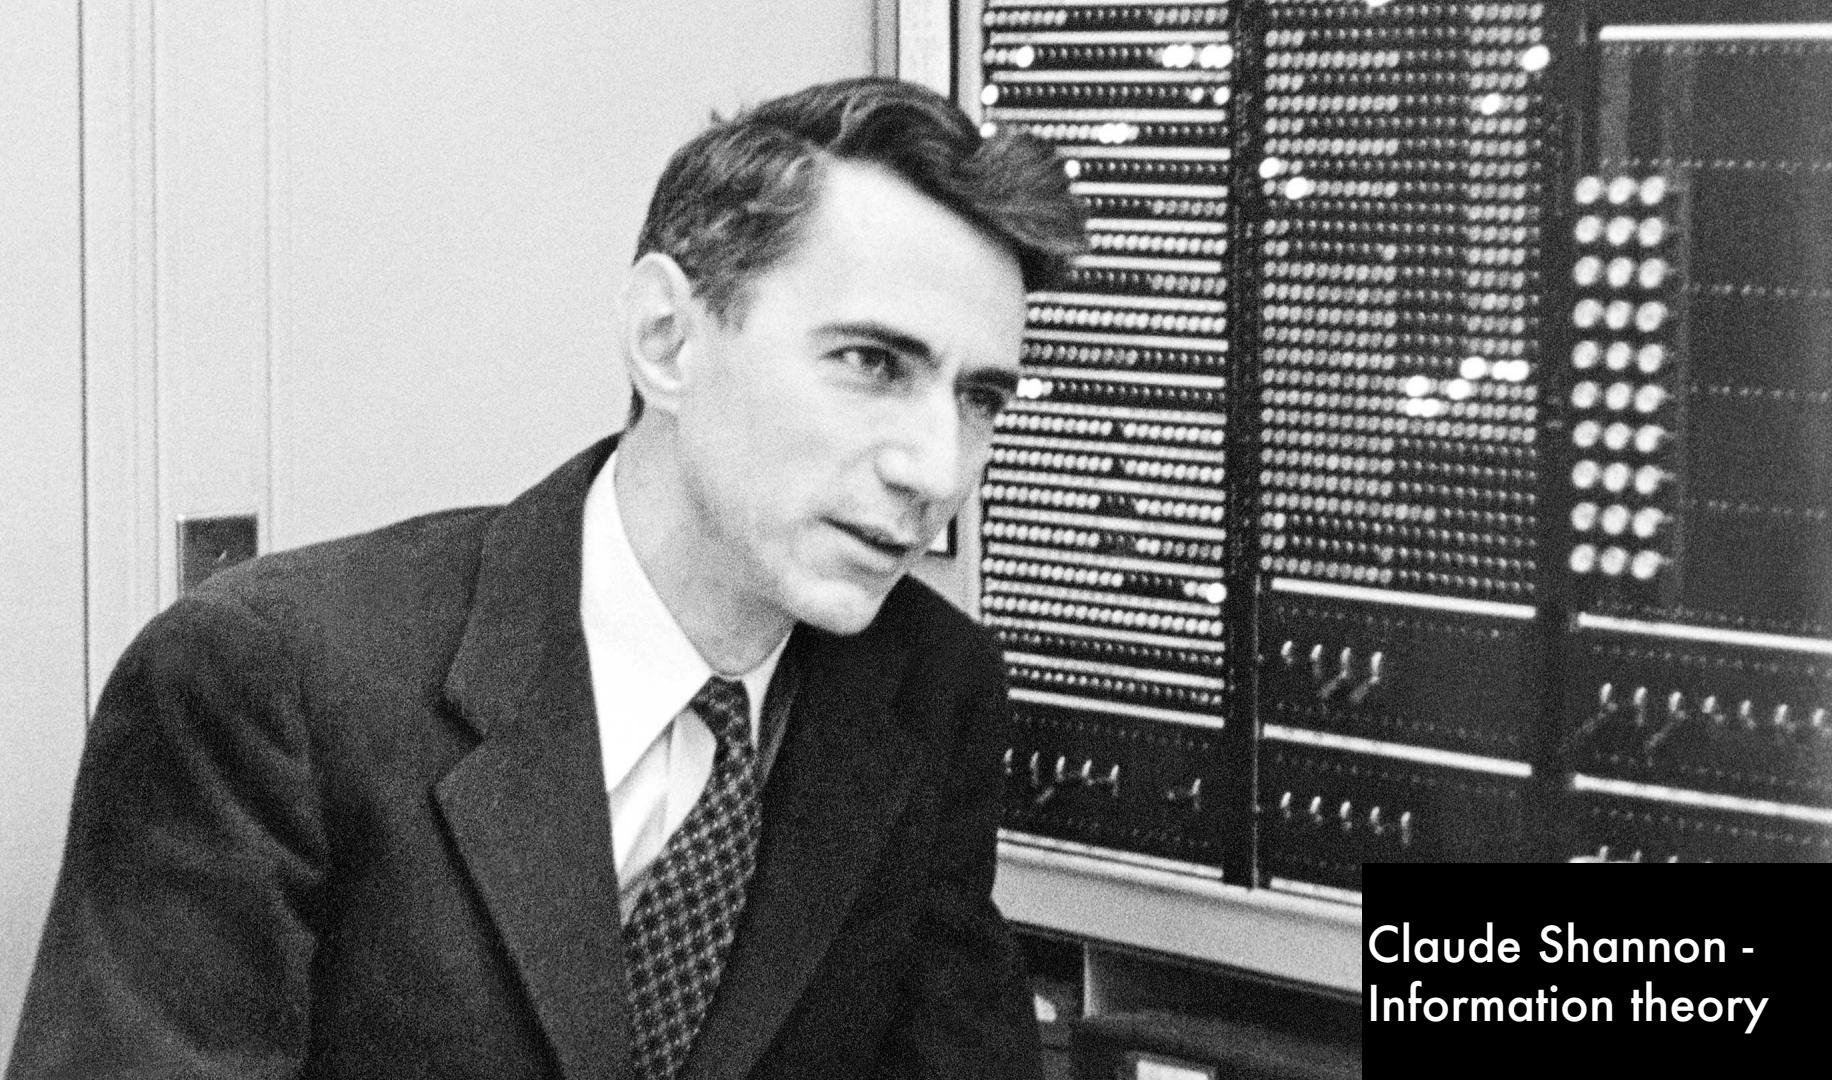
\$25.04

8 Used from \$21.06

From the bestselling author of the acclaimed *Chaos* and *Genius* comes a thoughtful and provocative exploration of the big ideas of the modern era: Information, communication, and information theory.

Acclaimed science writer James Gleick presents an eye-opening vision of how our relationship to information has transformed the very nature of human consciousness. A fascinating intellectual journey through the history of communication and information, from the language of Africa's talking drums to the invention of written alphabets; from the electronic transmission of code to the

[▼ Read more](#)



Claude Shannon -
Information theory

Before bits

- By 1948 more than 125 million conversations passed daily through the Bell System's 138 million miles of cable and 31 million telephone sets.
- Count words, characters, electricity consumption, minutes.

WHY BITS?

Physics! Electronic Implementation

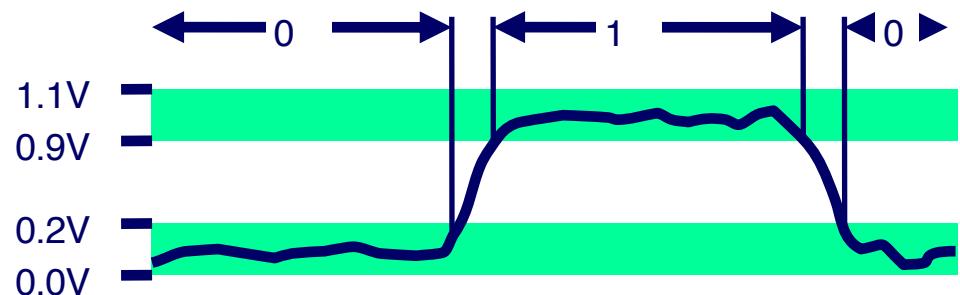
Easy to store with bistable elements.

e.g., high-low/off-on electromagnetism on disk.

Reliably transmitted on noisy and inaccurate wires

Very expressive and efficient.

Quantum computers?



COUNT EVERYTHING IN BINARY

Base 2 Number Representation

0, 1, 10, 11, 100, 101, ...

Represent 15213_{10} as 0011 1011 0110 1101₂

Represent 1.20_{10} as 1.0011 0011 0011 0011 [0011]...₂

Represent $(1.5213 \times 10^4)_{10}$ as $(1.1101 1011 0110 1 \times 2^{13})_2$

Represent negative numbers as ...?

(we'll come back to this)

Byte = 8 bits. Why?

Decimal: 0₁₀ to 255₁₀

$$255 = 2^8 - 1$$

Binary: 0000 0000₂ to 1111 1111₂

Hexadecimal: 00₁₆ to FF₁₆.

Why? Example: #FF5733, RGB 255, 87, 51

Base 16 number representation

Use characters '0' to '9' and 'A' to 'F'

Write in C with leading '0x', either case

$$0101\ 1010_2 = 0x5a = 0x5A = 0X5a$$

Hex Deci Biary

	Hex	Deci	Biary
0	0	0	0000
1	1	1	0001
2	2	2	0010
3	3	3	0011
4	4	4	0100
5	5	5	0101
6	6	6	0110
7	7	7	0111
8	8	8	1000
9	9	9	1001
A	10	10	1010
B	11	11	1011
C	12	12	1100
D	13	13	1101
E	14	14	1110
F	15	15	1111

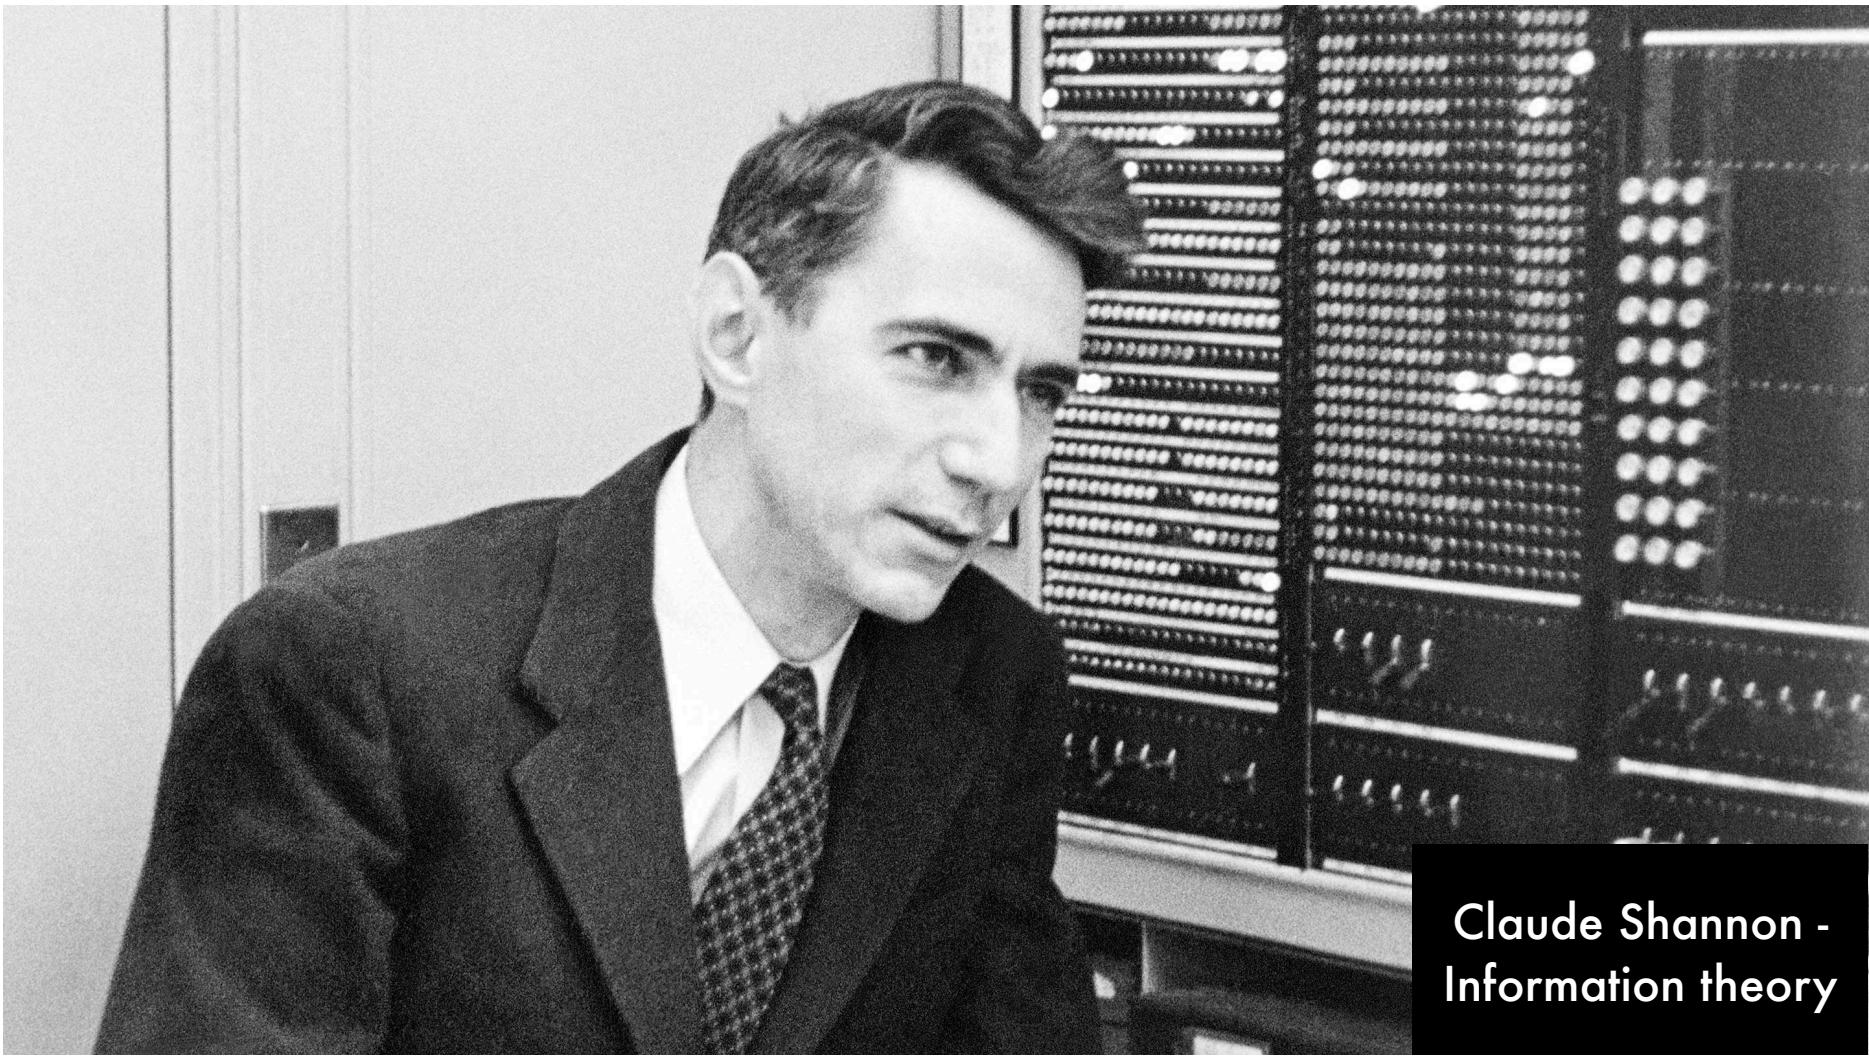
ENCODING BYTE VALUES

15213: 0011 1011 0110 1101

{ { { {
3 B 6 D



Name	Size ▾	Kind
HB50 cupcakes.JPG	2 MB	JPEG image
Roller Skating.JPG	1.3 MB	JPEG image
50HBJukebox2.jpg	720 KB	JPEG image
Facebook.tiff	399 KB	TIFF image
7_days_to_enrol.png	173 KB	PNG image
JoggingShoes.jpg	71 KB	JPEG image



Claude Shannon -
Information theory

Before bits

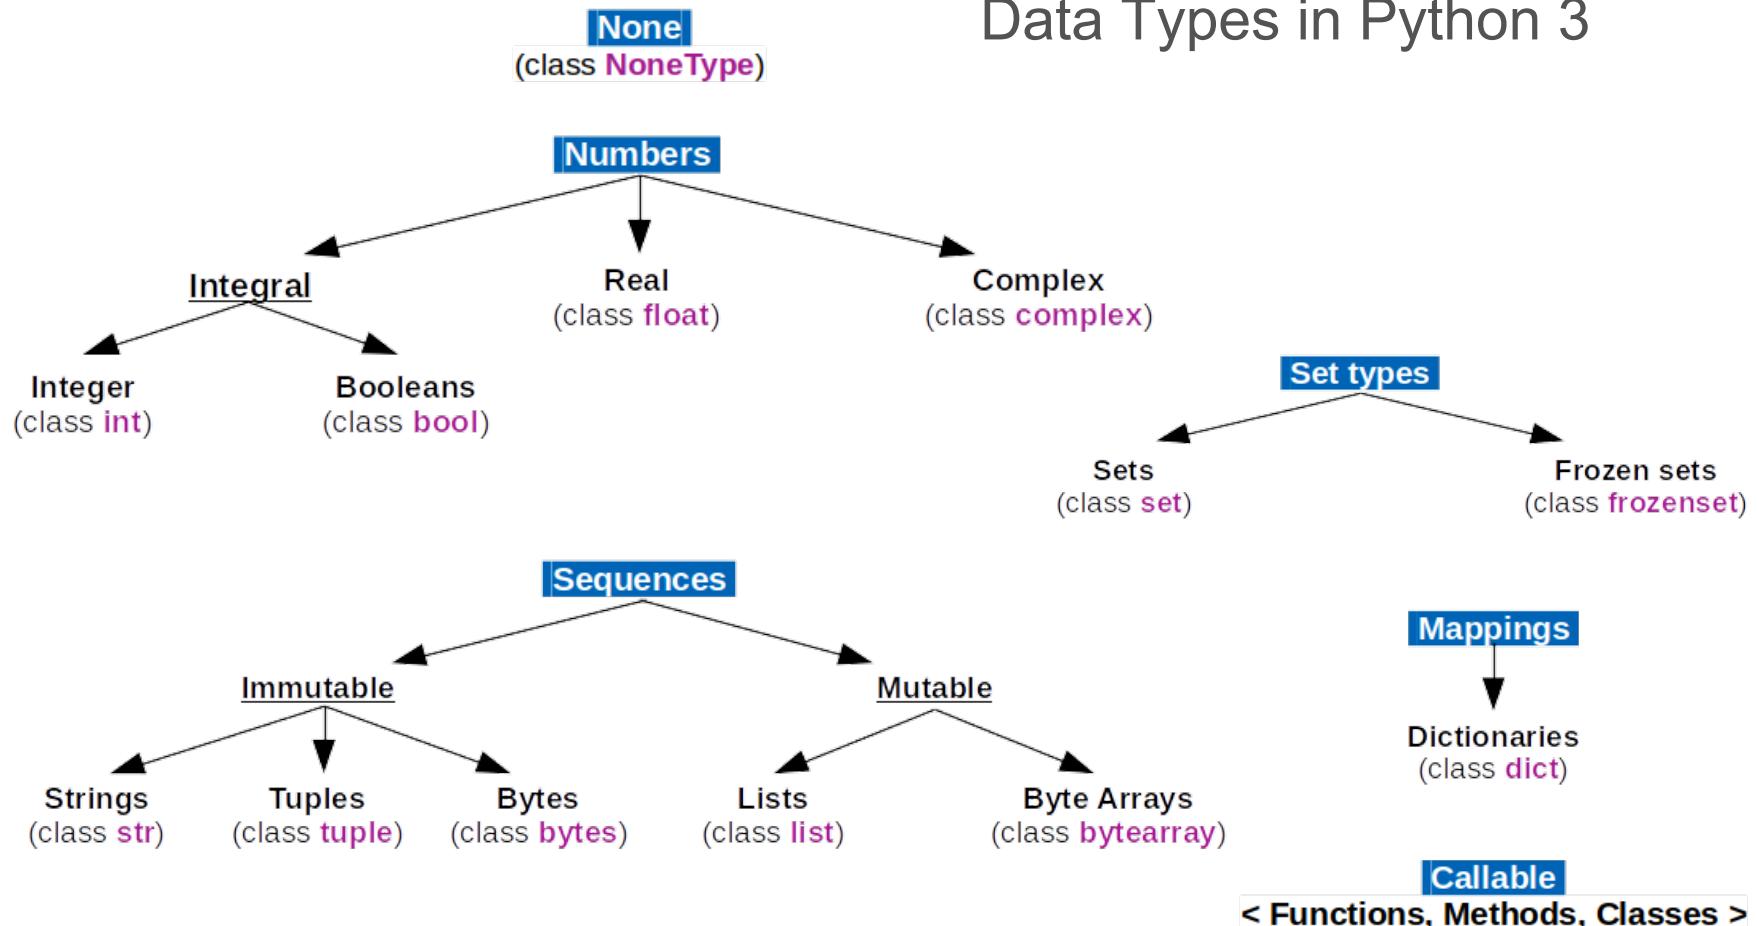
- By 1948 more than 125 million conversations passed daily through the Bell System's 138 million miles of cable and 31 million telephone sets.
- Count words, characters, electricity consumption, minutes.

Digital Representation of Data

- ❖ **Bits:** All digital data are sequences of 0 & 1 (binary digits)
 - ❖ Amenable to high-low/off-on electromagnetism
 - ❖ Layers of *abstraction* to interpret bit sequences
- ❖ **Data type:** First layer of abstraction to interpret a bit sequence with a human-understandable category of information; interpretation fixed by the PL
 - ❖ Example common datatypes: Boolean, Byte, Integer, “floating point” number (Float), Character, and String
- ❖ **Data structure:** A second layer of abstraction to *organize* multiple instances of same or varied data types as a more complex object with specified properties
 - ❖ Examples: Array, Linked list, Tuple, Graph, etc.

Digital Representation of Data

Data Types in Python 3



Digital Representation of Data

- ❖ The size and *interpretation* of a data type depends on PL
- ❖ A **Byte** (B; 8 bits) is typically the basic unit of data types
- ❖ **Boolean:**
 - ❖ Examples in data sci.: Y/N or T/F responses
 - ❖ Just 1 bit needed but actual size is almost always 1B, i.e., 7 bits are wasted! (**Q: Why?**)
- ❖ **Integer:**
 - ❖ Examples in data science: #friends, age, #likes
 - ❖ Typically 4 bytes; many variants (short, unsigned, etc.)
 - ❖ Java *int* can represent -2^{31} to $(2^{31} - 1)$; C *unsigned int* can represent 0 to $(2^{32} - 1)$; Python3 *int* is effectively unlimited length (PL magic!)

Digital Representation of Data

Q: *How many unique data items can be represented by 3 bytes?*

- ❖ Given k bits, we can represent 2^k unique data items
- ❖ 3 bytes = 24 bits $\Rightarrow 2^{24}$ items, i.e., 16,777,216 items
- ❖ Common approximation: 2^{10} (i.e., 1024) $\sim 10^3$ (i.e., 1000); recall kibibyte (KiB) vs kilobyte (KB) and so on

Q: *How many bits are needed to distinguish 97 data items?*

- ❖ For k unique items, invert the exponent to get $\log_2(k)$
- ❖ But #bits is an integer! So, we only need $\lceil \log_2(k) \rceil$
- ❖ So, we only need the next higher power of 2
- ❖ 97 $\rightarrow 128 = 2^7$; so, 7 bits

Digital Representation of Data

Q: How to convert from decimal to binary representation?

- Given decimal n, if power of 2 (say, 2^k), put 1 at bit position k; if $k=0$, stop; else pad with trailing 0s till position 0
- If n is not power of 2, identify the power of 2 just below n (say, 2^k); #bits is then k; put 1 at position k
- Reset n as $n - 2^k$; return to Steps 1-2
- Fill remaining positions in between with 0s

	7	6	5	4	3	2	1	0	Position/Exponent of 2
Decimal	128	64	32	16	8	4	2	1	Power of 2
5_{10}						1	0	1	
47_{10}				1	0	1	1	1	
163_{10}	1	0	1	0	0	0	1	1	
16_{10}				1	0	0	0	0	

Q: Binary to decimal?

Digital Representation of Data

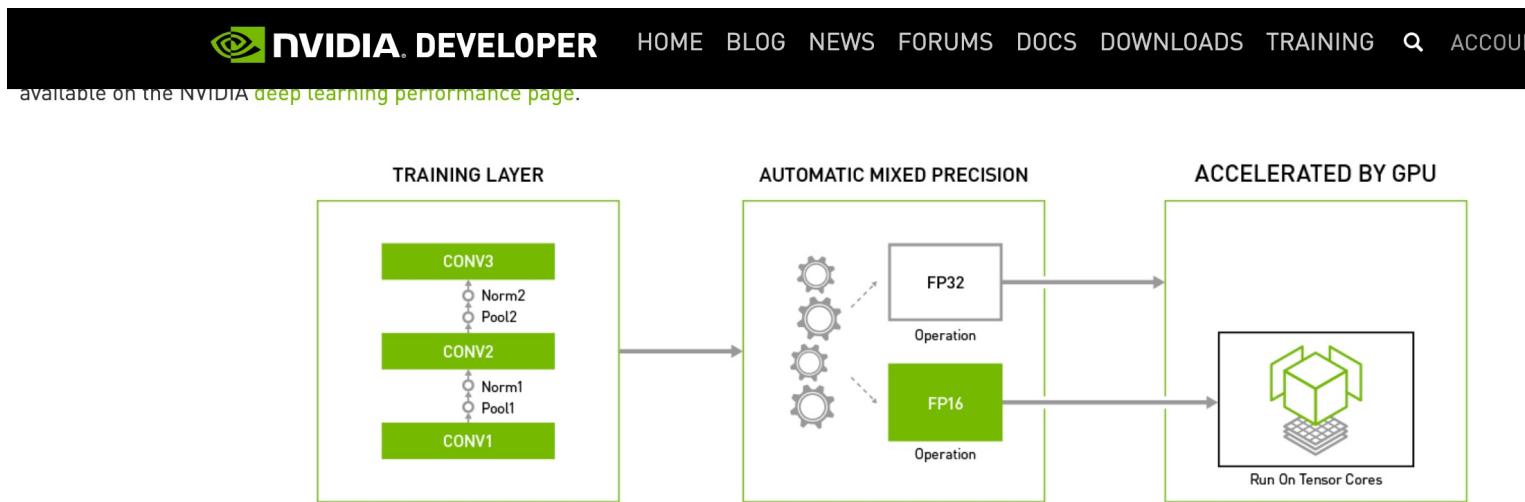
- ❖ *Hexadecimal* representation is a common stand-in for binary representation; more succinct and readable
 - ❖ Base 16 instead of base 2 cuts display length by ~4x
 - ❖ Digits are 0, 1, ... 9, A (10_{10}), B, ... F (15_{10})
 - ❖ From binary: combine 4 bits at a time from lowest

Decimal	Binary	Hexadecimal	
5_{10}	101_2	5_{16}	Alternative notations
47_{10}	$10\ 1111_2$	$2F_{16}$	
163_{10}	$1010\ 0011_2$	$A3_{16}$	$0xA3$ or $A3H$
16_{10}	$1\ 0000_2$	$1\ 0_{16}$	

Digital Representation of Data

- ❖ **Float:**

- ❖ Examples in data sci.: salary, scores, model weights
- ❖ IEEE-754 single-precision format is 4B long; double-precision format is 8B long
- ❖ Java and C *float* is single; Python *float* is double!

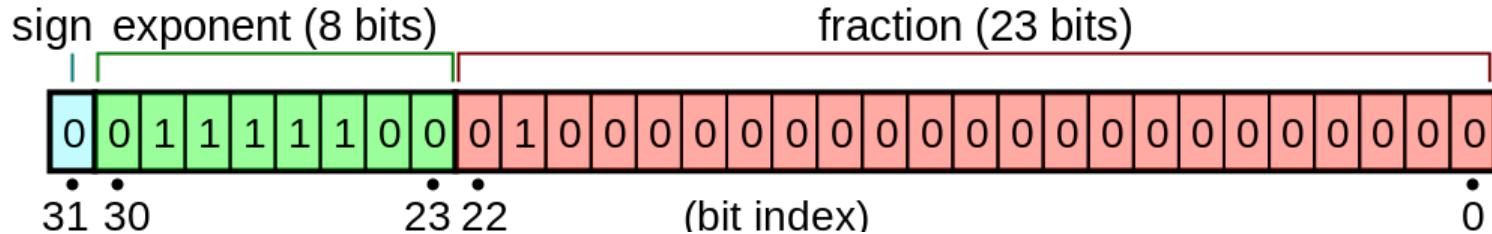


Using Automatic Mixed Precision for Major Deep Learning Frameworks

Digital Representation of Data

- ❖ **Float:**

- ❖ Standard IEEE format for single (aka binary32):



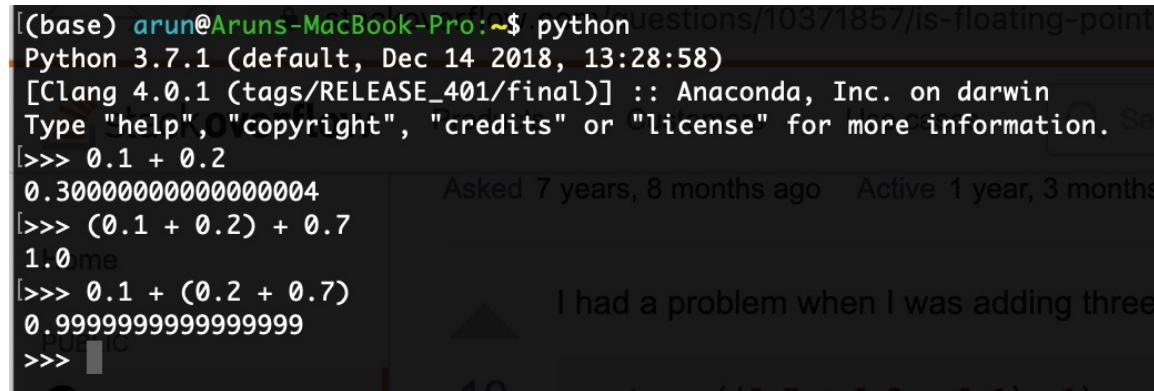
$$(-1)^{sign} \times 2^{exponent-127} \times \left(1 + \sum_{i=1}^{23} b_{23-i} 2^{-i}\right)$$

$$(-1)^0 \times 2^{124-127} \times (1 + 1 \cdot 2^{-2}) = (1/8) \times (1 + (1/4)) = 0.15625$$

(NB: Converting decimal reals/fractions to float is NOT in syllabus!) 57

Digital Representation of Data

- ❖ Due to representation imprecision issues, floating point arithmetic (addition and multiplication) is not associative!



A screenshot of a Stack Overflow question titled "I had a problem when I was adding three floating point numbers". The question asks why the sum of three floating-point numbers does not add up to exactly 1.0. The code shown in the screenshot is:

```
[base] arun@Arun's-MacBook-Pro:~$ python
Python 3.7.1 (default, Dec 14 2018, 13:28:58)
[Clang 4.0.1 (tags/RELEASE_401/final)] :: Anaconda, Inc. on darwin
Type "help", "copyright", "credits" or "license" for more information.
[>>> 0.1 + 0.2
0.30000000000000004
[>>> (0.1 + 0.2) + 0.7
1.0
[>>> 0.1 + (0.2 + 0.7)
0.9999999999999999
>>>
```

The question has been asked 7 years, 8 months ago and is active 1 year, 3 months.

- ❖ In binary32, special encodings recognized:
 - ❖ Exponent 0xFF and fraction 0 is +/- “Infinity”
 - ❖ Exponent 0xFF and fraction <> 0 is “NaN”
 - ❖ Max is $\sim 3.4 \times 10^{38}$; min +ve is $\sim 1.4 \times 10^{-45}$

Digital Representation of Data

- ❖ More float standards: double-precision (float64; 8B) and half-precision (float16; 2B); different #bits for exponent, fraction
- ❖ Float16 is now common for *deep learning* parameters:
 - ❖ Native support in PyTorch, TensorFlow, etc.; APIs also exist for weight quantization/rounding post training
 - ❖ NVIDIA Deep Learning SDK support mixed-precision training; 2-3x speedup with similar accuracy!
- ❖ New processor hardware (FPGAs, ASICs, etc.) enable arbitrary precision, even 1-bit (!), but accuracy is lower