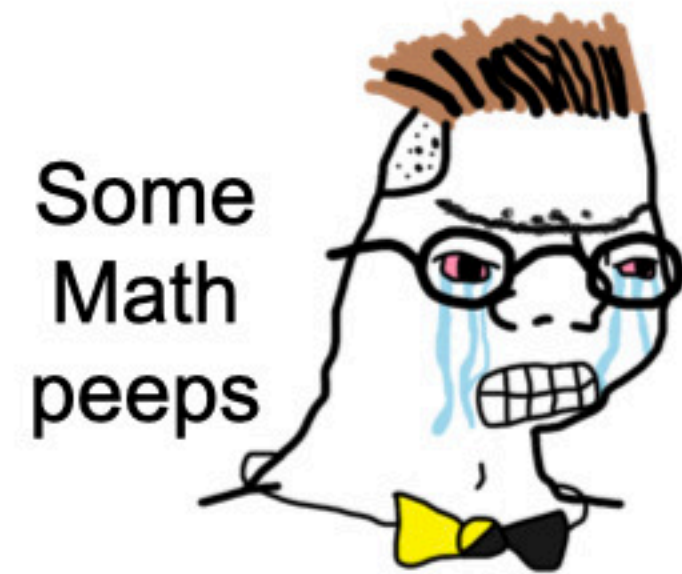


THEN (1980s)



Noooo... CS is just applied math!



Noooo... CS is just applied EE!

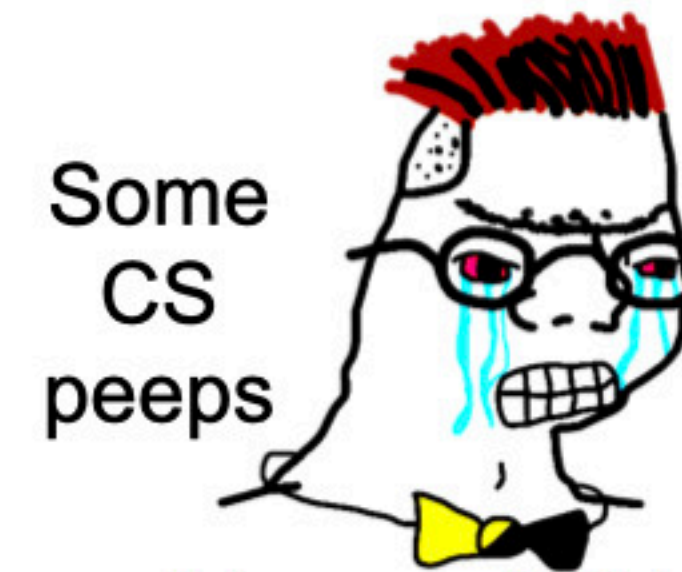


Haha. Computer Science go brrr

NOW (2020s)



Noooo... DS is just applied math!



Noooo... DS is just applied CS!



Haha. Data Science go brrr

DSC 102

Systems for Scalable Analytics

- Haojian Jin

Where are we in the class?

- ❖ Basics of Computer Organization
 - ❖ Digital Representation of Data
 - ❖ **Processors and Memory Hierarchy**
- ❖ Basics of Operating Systems
 - ❖ Process Management: Virtualization; Concurrency
 - ❖ Filesystem and Data Files
 - ❖ Main Memory Management
- ❖ Persistent Data Storage

Today

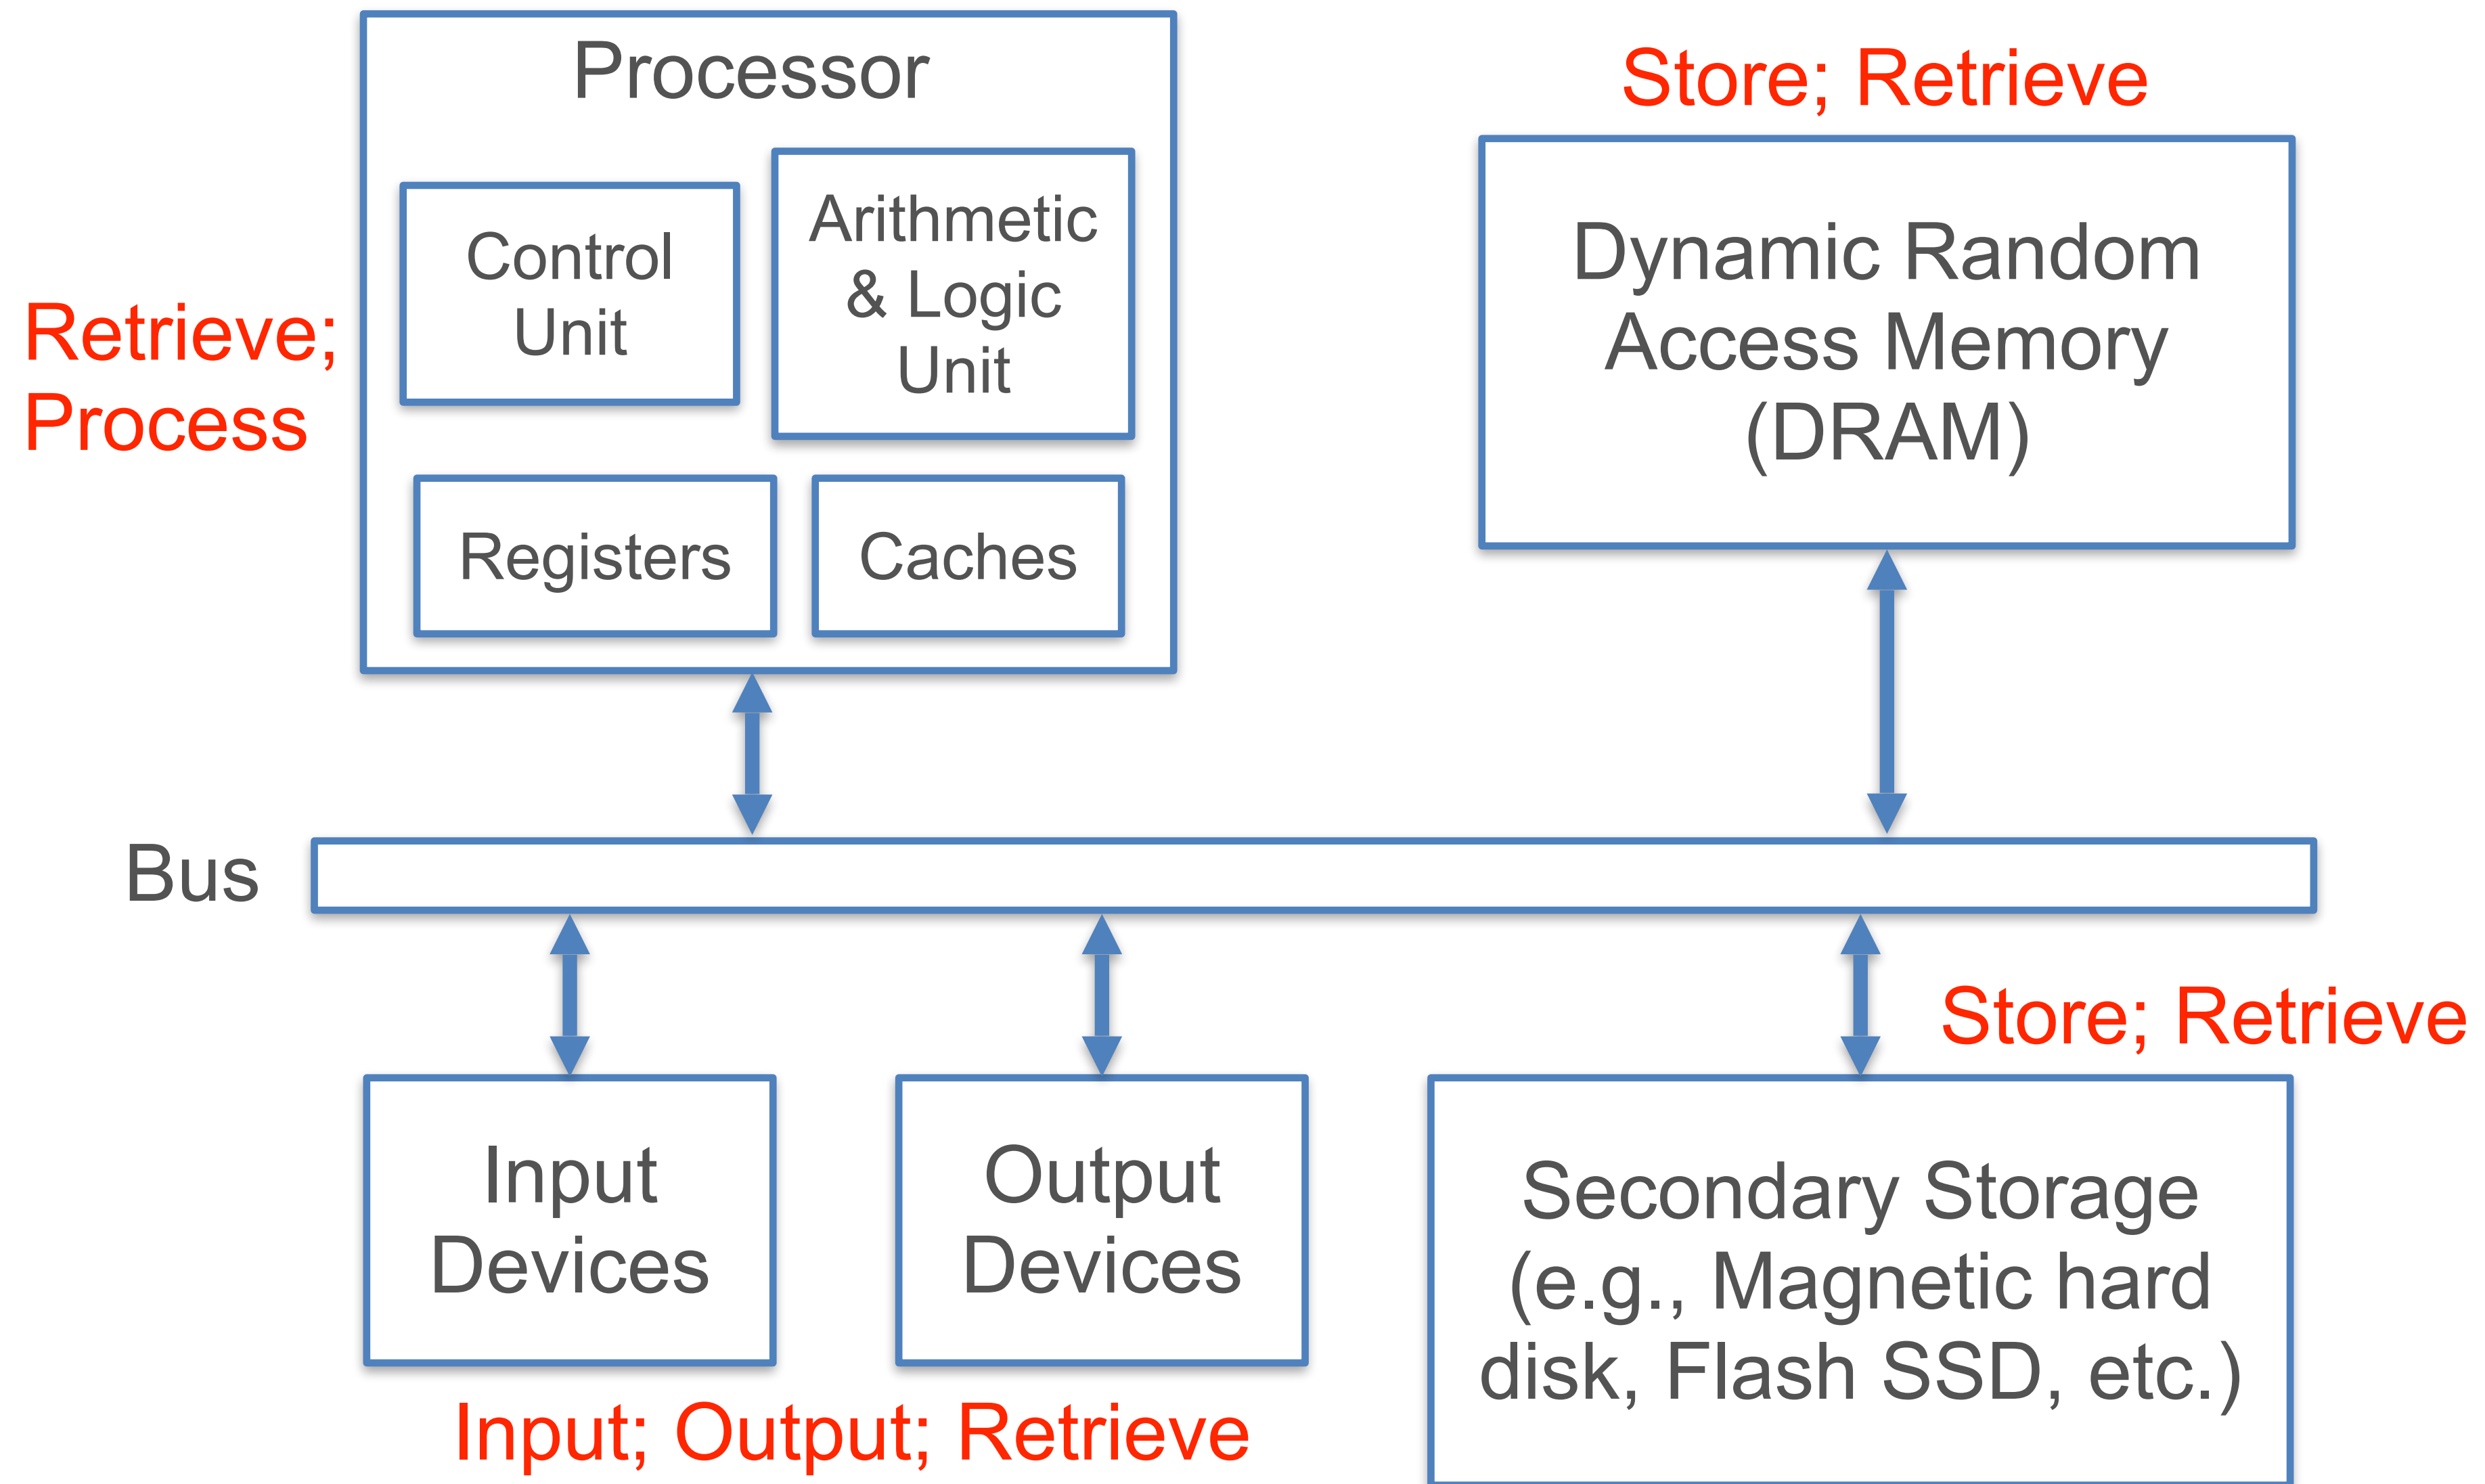
The memory abstraction

Locality of reference

The memory hierarchy


Storage technologies and trends

Abstract Computer Parts and Data



Memory Hierarchy in PA0

- Pandas DataFrame needs data to fit entirely in DRAM
- **Dask DataFrame** automatically manages Disk vs DRAM for you
 - Full data sits on Disk, brought to DRAM upon compute()
 - Dask stages out computations using Pandas



- **Tradeoff:** Dask may throw memory configuration issues. :)

Instruction

Register names



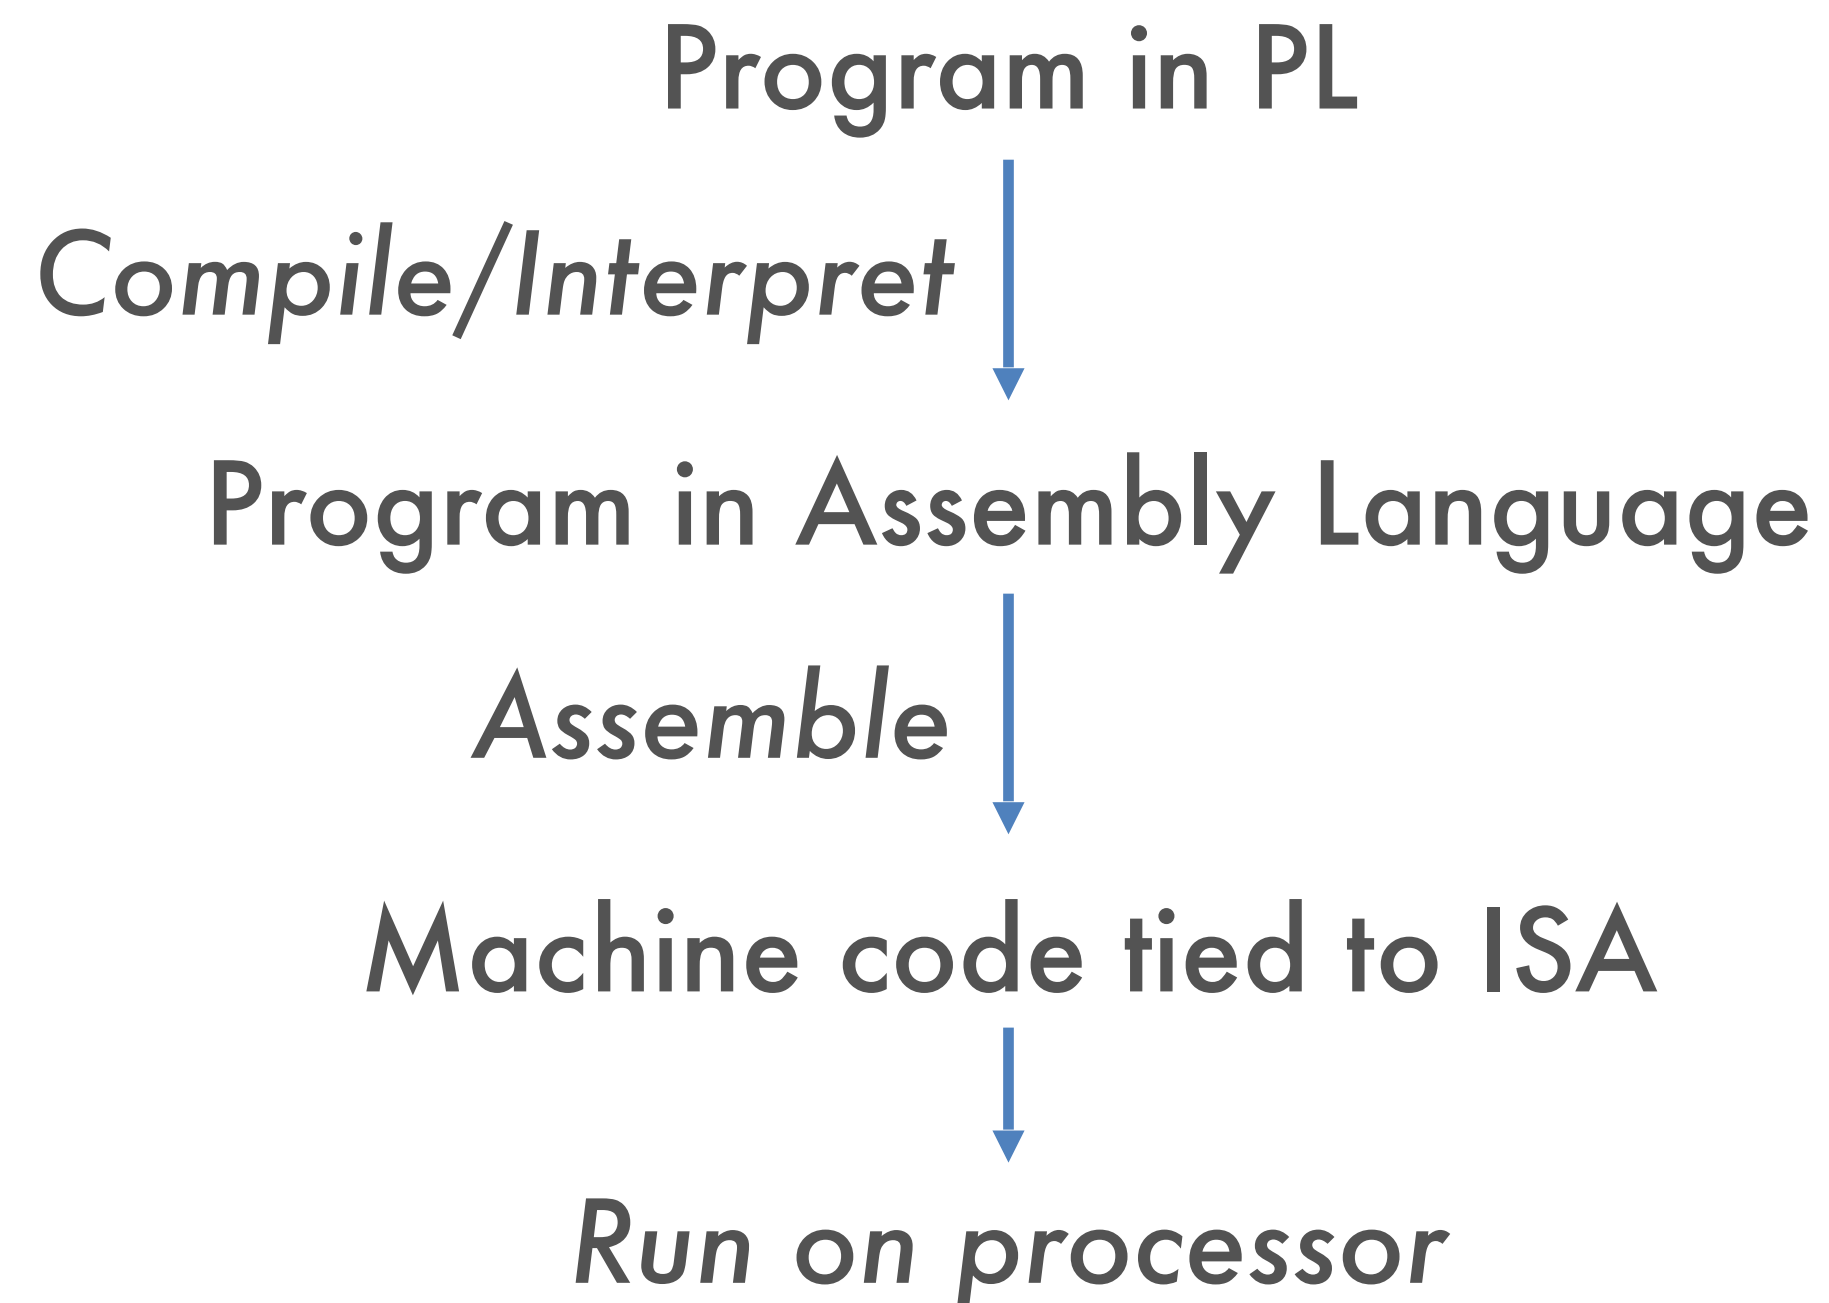
addq %rbx, %rax

is

rax += rbx

Basics of Processors

- **Processor:** Hardware to orchestrate and execute *instructions* to *manipulate data* as specified by a program
 - Examples: CPU, GPU, FPGA, TPU, embedded, etc.
- **ISA (Instruction Set Architecture):**
 - The vocabulary of commands of a processor



```
80483b4: 55          push    %ebp
80483b5: 89 e5       mov     %esp,%ebp
80483b7: 83 e4 f0    and     $0xffffffff0,%esp
80483ba: 83 ec 20    sub     $0x20,%esp
80483bd: c7 44 24 1c 00 00 00 movl    $0x0,0x1c(%esp)
80483c4: 00
80483c5: eb 11       jmp     80483d8 <main+0x24>
80483c7: c7 04 24 b0 84 04 08 movl    $0x80484b0,(%esp)
80483ce: e8 1d ff ff ff call    80482f0 <puts@plt>
80483d3: 83 44 24 1c 01 addl    $0x1,0x1c(%esp)
80483d8: 83 7c 24 1c 09 cmpl    $0x9,0x1c(%esp)
80483dd: 7e e8       jle     80483c7 <main+0x13>
80483df: b8 00 00 00 00 mov     $0x0,%eax
80483e4: c9         leave
80483e5: c3         ret
80483e6: 90         nop
80483e7: 90         nop
80483e8: 90         nop
80483e9: 90         nop
80483ea: 90         nop
```

Basics of Processors

Q: *How does a processor execute machine code?*

- Most common approach: **load-store architecture**
- **Registers:** Tiny local memory (“scratch space”) on proc. into which instructions and data are copied
- ISA specifies bit length/format of machine code commands
- ISA has several commands to manipulate register contents

Writing & Reading Memory

Write

Transfer data from memory to CPU

movq %rax, %rsp

“Store” operation


Read

Transfer data from CPU to memory

movq %rsp, %rax


“Load” operation

Abstract Computer Parts and Data

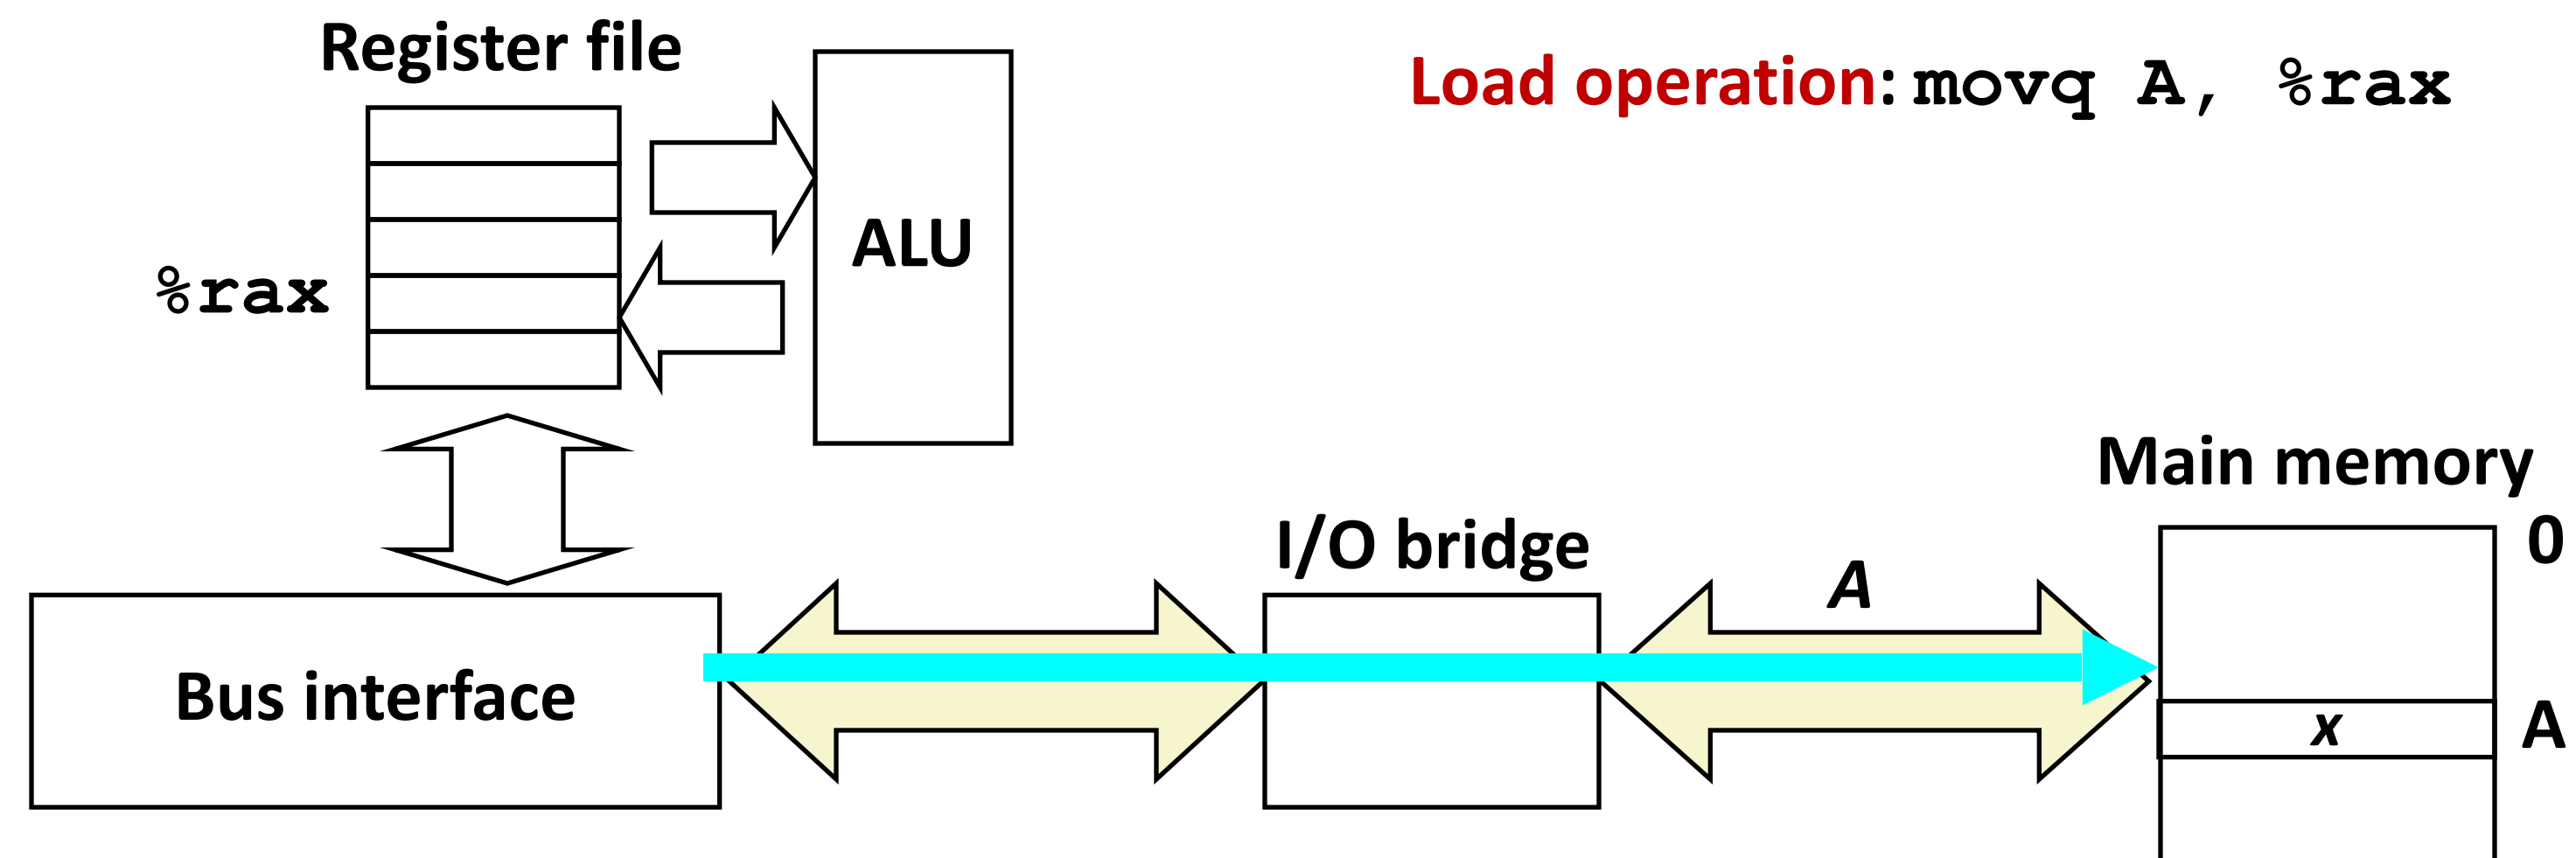


Bus Structure Connecting CPU and Memory

A **bus** is a collection of parallel wires that carry address, data, and control signals. Buses are typically shared by multiple devices.




Memory Read Transaction (1)




CPU places address *A* on the memory bus.

Memory Read Transaction (2)




Main memory reads `A` from the memory bus, retrieves word `x`, and places it on the bus.

Memory Read Transaction (3)




CPU reads word `x` from the bus and copies it into register `%rax`.

Memory Write Transaction (1)




CPU places address `A` on bus. Main memory reads it and waits for the corresponding data word to arrive.

Memory Write Transaction (2)



CPU places data word `y` on the bus.

Memory Write Transaction (3)



Main memory reads data word `y` from the bus and stores it at address `A`.

Basics of Processors

Q: *How does a processor execute machine code?*

- Types of ISA commands to manipulate register contents:
 - **Memory access: load** (copy bytes from a DRAM address to register); **store** (reverse); put constant
 - **Arithmetic & logic** on data items in registers: add/multiply/etc.; bitwise ops; compare, etc.; handled by ALU
 - **Control flow** (branch, call, etc.); handled by CU
- **Caches:** Small local memory to buffer instructions/data

80483b5:	89 e5	mov	%esp,%ebp
80483b7:	83 e4 f0	and	\$0xfffffffff0,%esp
80483ba:	83 ec 20	sub	\$0x20,%esp
80483bd:	c7 44 24 1c 00 00 00	movl	\$0x0,0x1c(%esp)
80483c4:	00		
80483c5:	eb 11	jmp	80483d8 <main+0x24>
80483c7:	c7 04 24 b0 84 04 08	movl	\$0x80484b0, (%esp)
80483ce:	e8 1d ff ff ff	call	80482f0 <puts@plt>
80483d3:	83 44 24 1c 01	addl	\$0x1,0x1c(%esp)
80483d8:	83 7c 24 1c 09	cmpl	\$0x9,0x1c(%esp)
80483dd:	7e e8	jle	80483c7 <main+0x13>
80483df:	b8 00 00 00 00	mov	\$0x0,%eax
80483e4:	c9	leave	
80483e5:	c3	ret	
80483e6:	90	nop	
80483e7:	90	nop	
80483e8:	90	nop	
80483e9:	90	nop	

Example

Q: *What does this program do when run with 'python'?
(Assume tmp.csv is in current working directory)*

tmp.py

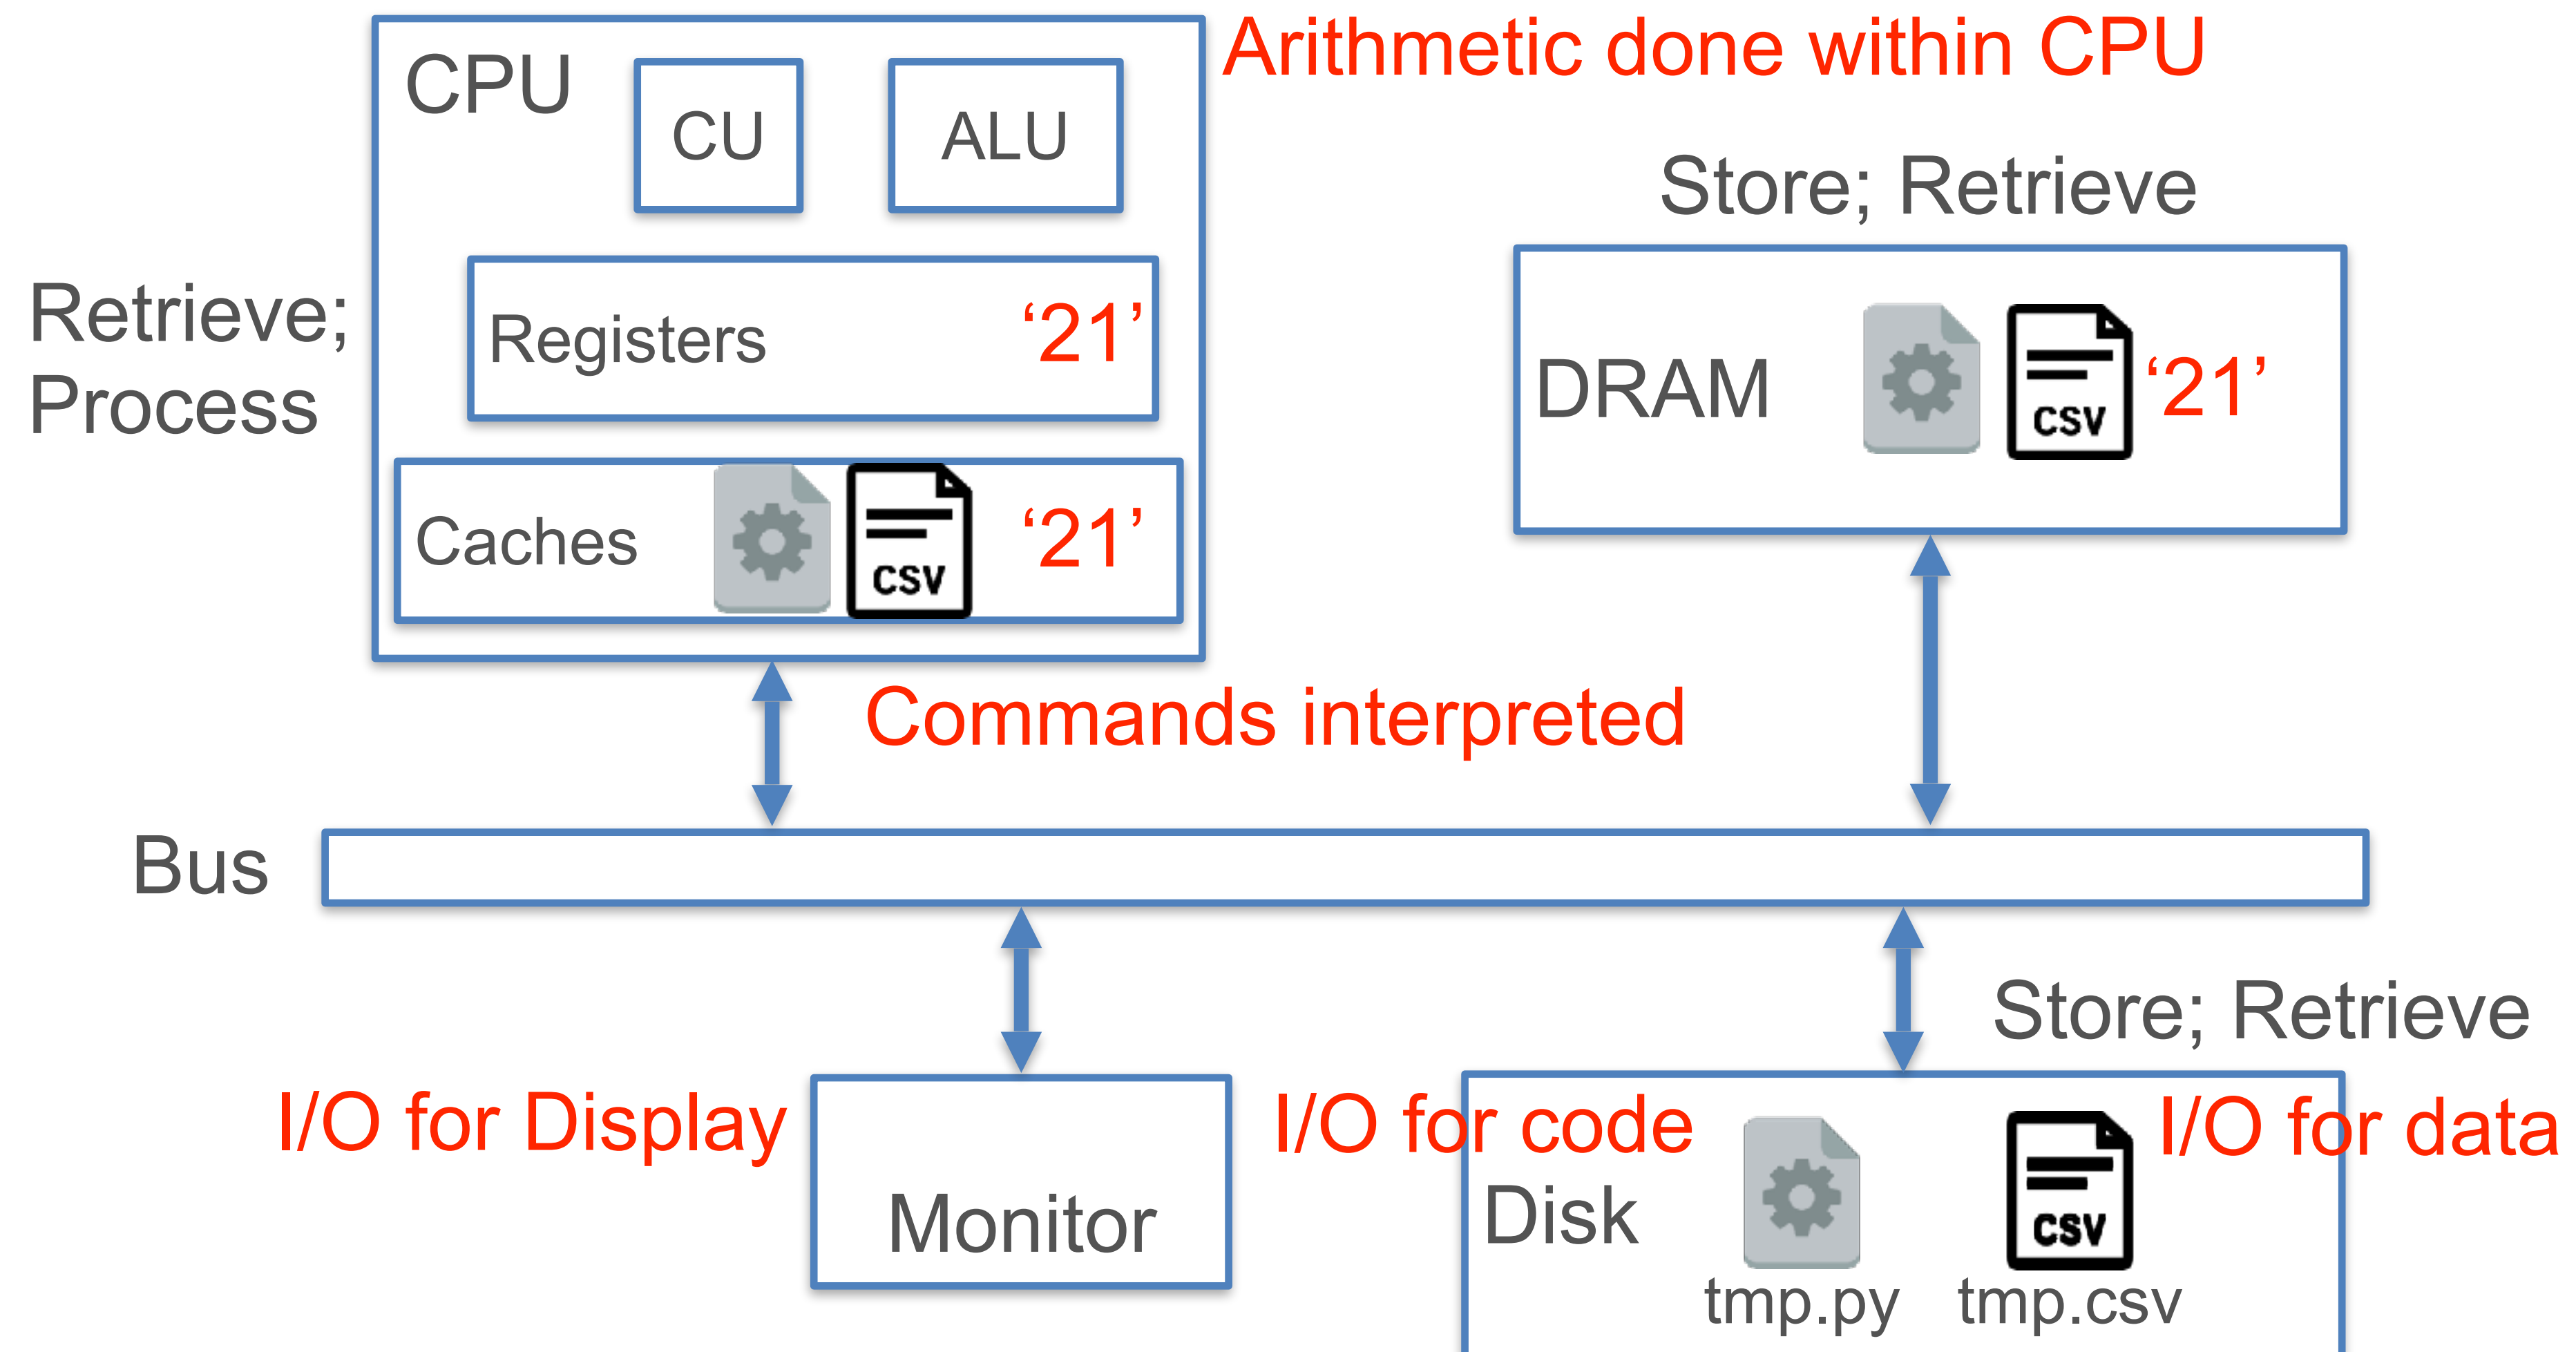
```
import pandas as p
m = p.read_csv('tmp.csv',header=None)
s = m.sum().sum()
print(s)
```

tmp.csv

```
1,2,3
4,5,6
```


Example

Rough sequence of events when program is executed



x86 v.s. x64

- x86 is the name of an ISA (80X86). 32-bit CPU.
- x64 is the name of an ISA(x86-64). 64-bit CPU
- 32-bit CPU can handle 32 bits information in each instruction.
- When we represent the memory address in bits, 32-bit CPU can support at most 2^{32} bytes = 4 GB.
- You can install a 32-bit OS on a 64-bit CPU. But not a 64-bit OS on a 32-bit CPU.

Today

The memory abstraction

Locality of reference

The memory hierarchy

Storage technologies and trends

copyij v.s copyji: copy a 2048 X 2048 integer array

```
void copyij(long int src[2048][2048], long int dst[2048][2048])
{
    long int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```


4.3 milliseconds

```
void copyji(long int src[2048][2048], long int dst[2048][2048])
{
    long int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8 milliseconds

The CPU-Memory Gap

The gap *widens* between DRAM, disk, and CPU speeds.



Locality to the Rescue!


The key to bridging this CPU-Memory gap is an important property of computer programs known as **locality**.

Locality

Principle of Locality: Many Programs tend to use data and instructions with addresses near or equal to those they have used recently.


Temporal locality:

Recently referenced items are likely to be referenced again in the near future



Spatial locality:

Items with nearby addresses tend to be referenced close together in time



Locality Example

```
num_list = [1, 2, 3, 4, 5, 7]
sum = 0;
for (x in num_list)
    sum += x;
return sum;
```

**Spatial or Temporal
Locality?**

Data references

Reference array elements in succession (stride-1 reference pattern).

Reference variable **sum** each iteration.

spatial

temporal

Instruction references

Reference instructions in sequence.

Cycle through loop repeatedly.

spatial

temporal

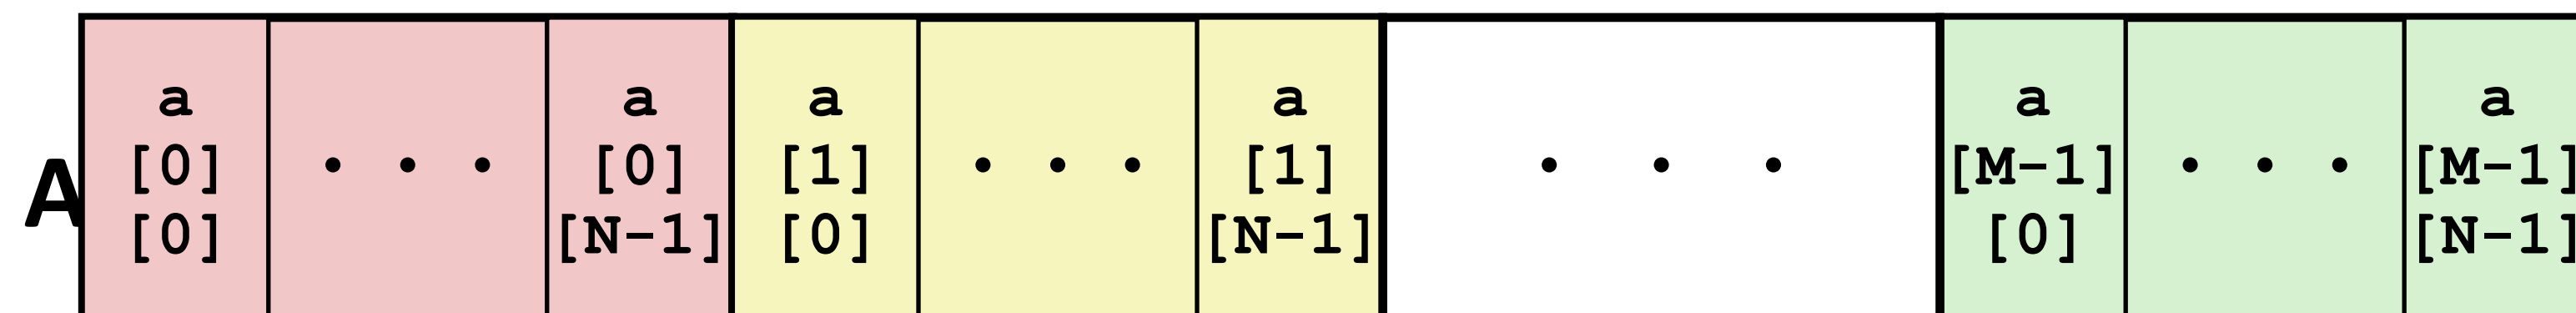
Qualitative Estimates of Locality

**Hint: array layout
is row-major order**

```
int sum_array_rows(int a[M][N])
{
    int i, j, sum = 0;

    for (i = 0; i < M; i++)
        for (j = 0; j < N; j++)
            sum += a[i][j];

    return sum;
}
```



Claim: Being able to look at code and get a qualitative sense of its locality is a key skill for a professional programmer.

Question: Does this function have good locality with respect to array *a*?

Locality Example

```
int sum_array_cols(int a[M][N])
{
    int i, j, sum = 0;


    for (j = 0; j < N; j++)
        for (i = 0; i < M; i++)
            sum += a[i][j];

    return sum;
}
```

Answer: no, unless...

M is very small

Question: Does this function have good locality with respect to array a?



Locality Example

```
int sum_array_3d(int a[M][N][N])
{
    int i, j, k, sum = 0;

    for (i = 0; i < N; i++)
        for (j = 0; j < N; j++)
            for (k = 0; k < M; k++)
                sum += a[k][i][j];

    return sum;
}
```

Question: Can you permute the loops so that the function scans the 3-d array *a* with a stride-1 reference pattern (and thus has good spatial locality)?

Answer: make j the inner loop

Today

The memory abstraction

Locality of reference

The memory hierarchy

Storage technologies and trends

Memory Hierarchies

- Some fundamental and enduring properties of hardware and software:
 - Fast storage technologies cost more per byte, have less capacity, and require more power (heat!).
 - The gap between CPU and main memory speed is widening.
 - Well-written programs tend to exhibit good locality.
- These properties complement each other well for many types of programs.
- They suggest an approach for organizing memory and storage systems known as a **memory hierarchy**.


Example Memory Hierarchy



Smaller,
faster,
and
costlier
(per byte)
storage
devices




Larger,
slower,
and
cheaper
(per byte)
storage
devices




Caches

- **Cache:** A smaller, faster storage device that acts as a staging area for a subset of the data in a larger, slower device.
- Fundamental idea of a memory hierarchy:
 - For each k , the faster, smaller device at level k serves as a cache for the larger, slower device at level $k+1$.
- Why do memory hierarchies work?
 - Because of locality: programs tend to access the data at level k more often than they access the data at level $k+1$.
 - Thus, the storage at level $k+1$ can be slower, and thus larger and cheaper per bit.
- **Big Idea (Ideal):** The memory hierarchy creates a large pool of storage that costs as much as the cheap storage near the bottom, but that serves data to programs at the rate of the fast storage near the top.


General Cache Concepts



General Cache Concepts: Hit



General Cache Concepts: Miss



Data in block b is needed

Block b is not in cache:
Miss!

Block b is fetched from memory

Block b is stored in cache

- **Placement policy:**
determines where b goes
- **Replacement policy:**
determines which block gets evicted (victim)

Examples of Caching in the Mem. Hierarchy

Cache Type	What is Cached?	Where is it Cached?	Latency (cycles)	Managed By
Registers	4-8 byte words	CPU core	0	Compiler
TLB	Address translations	On-Chip TLB	0	Hardware MMU
L1 cache	64-byte blocks	On-Chip L1	4	Hardware
L2 cache	64-byte blocks	On-Chip L2	10	Hardware
Virtual Memory	4-KB pages	Main memory	100	Hardware + OS
Buffer cache	Parts of files	Main memory	100	OS
Disk cache	Disk sectors	Disk controller	100,000	Disk firmware
Network buffer cache	Parts of files	Local disk	10,000,000	NFS client
Browser cache	Web pages	Local disk	10,000,000	Web browser
Web cache	Web pages	Remote server disks	1,000,000,000	Web proxy server

Recap

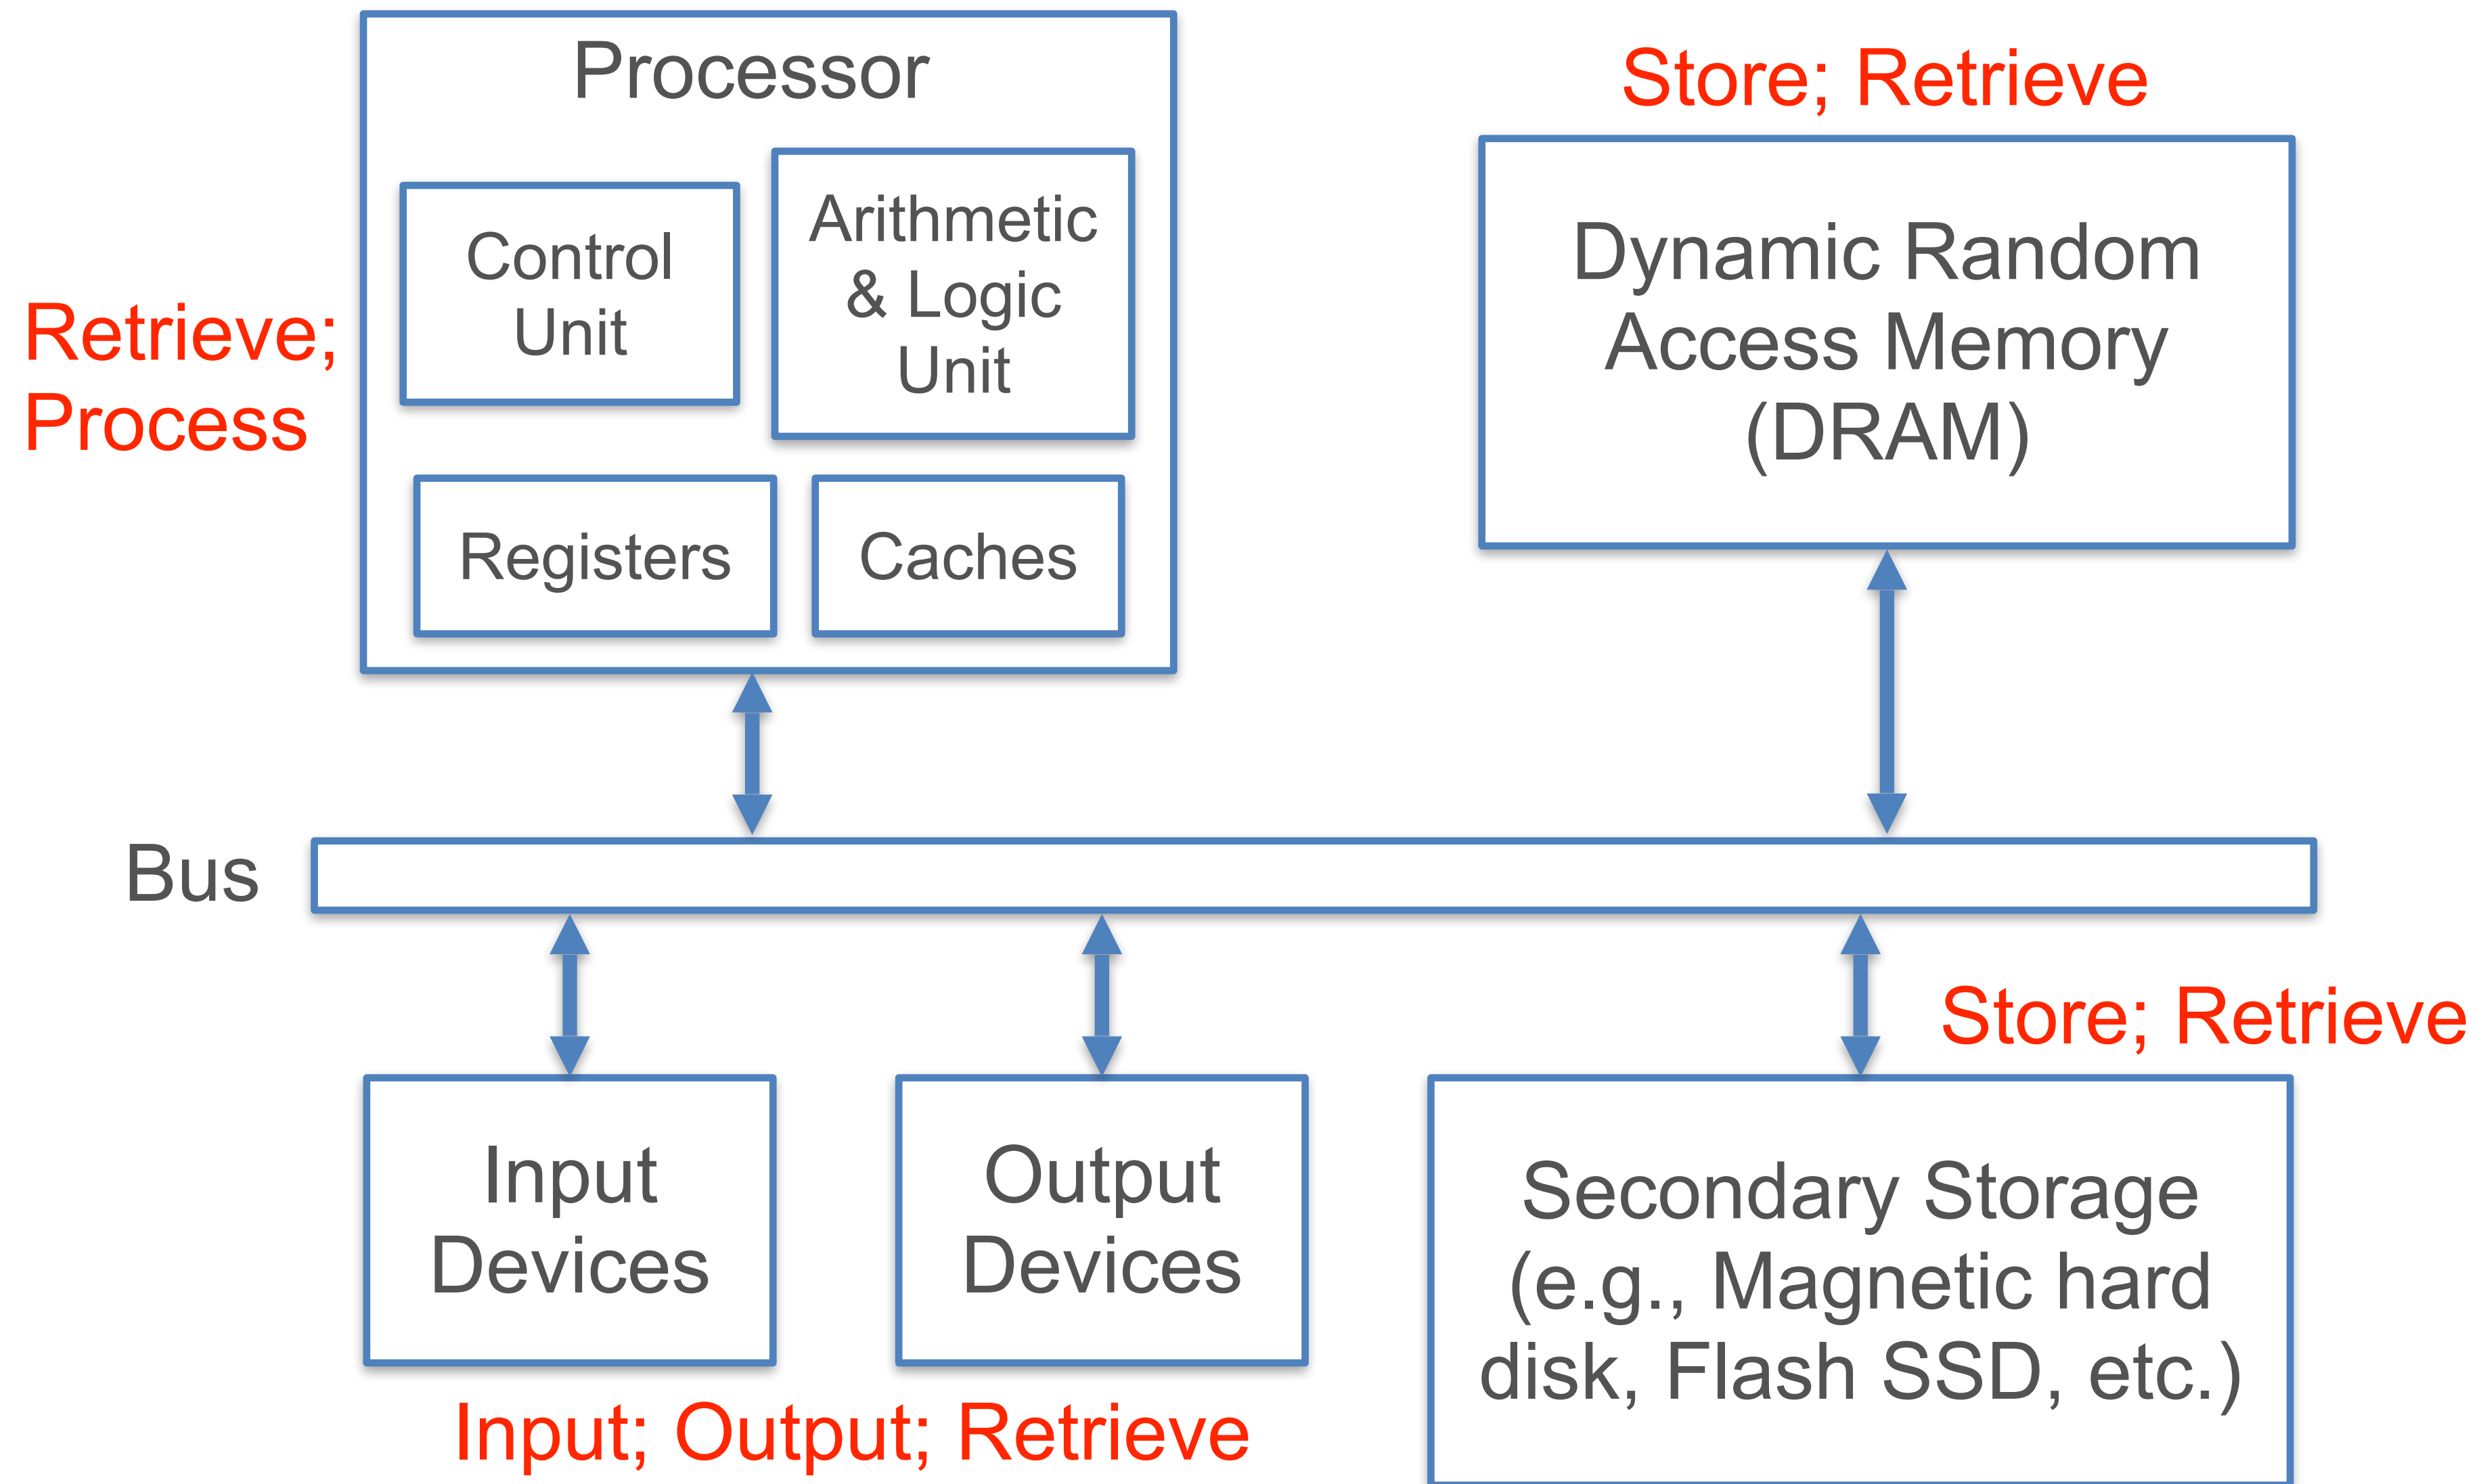
The memory abstraction

Locality of reference

The memory hierarchy

Storage technologies and trends

Abstract Computer Parts and Data



copyij v.s copyji: copy a 2048 X 2048 integer array

```
void copyij(long int src[2048][2048], long int dst[2048][2048])
{
    long int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

4.3 milliseconds

```
void copyji(long int src[2048][2048], long int dst[2048][2048])
{
    long int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

81.8 milliseconds


Example Memory Hierarchy



Smaller,
faster,
and
costlier
(per byte)
storage
devices



Larger,
slower,
and
cheaper
(per byte)
storage
devices



Today

The memory abstraction

Locality of reference

The memory hierarchy

Storage technologies and trends

What's Inside A Disk Drive?

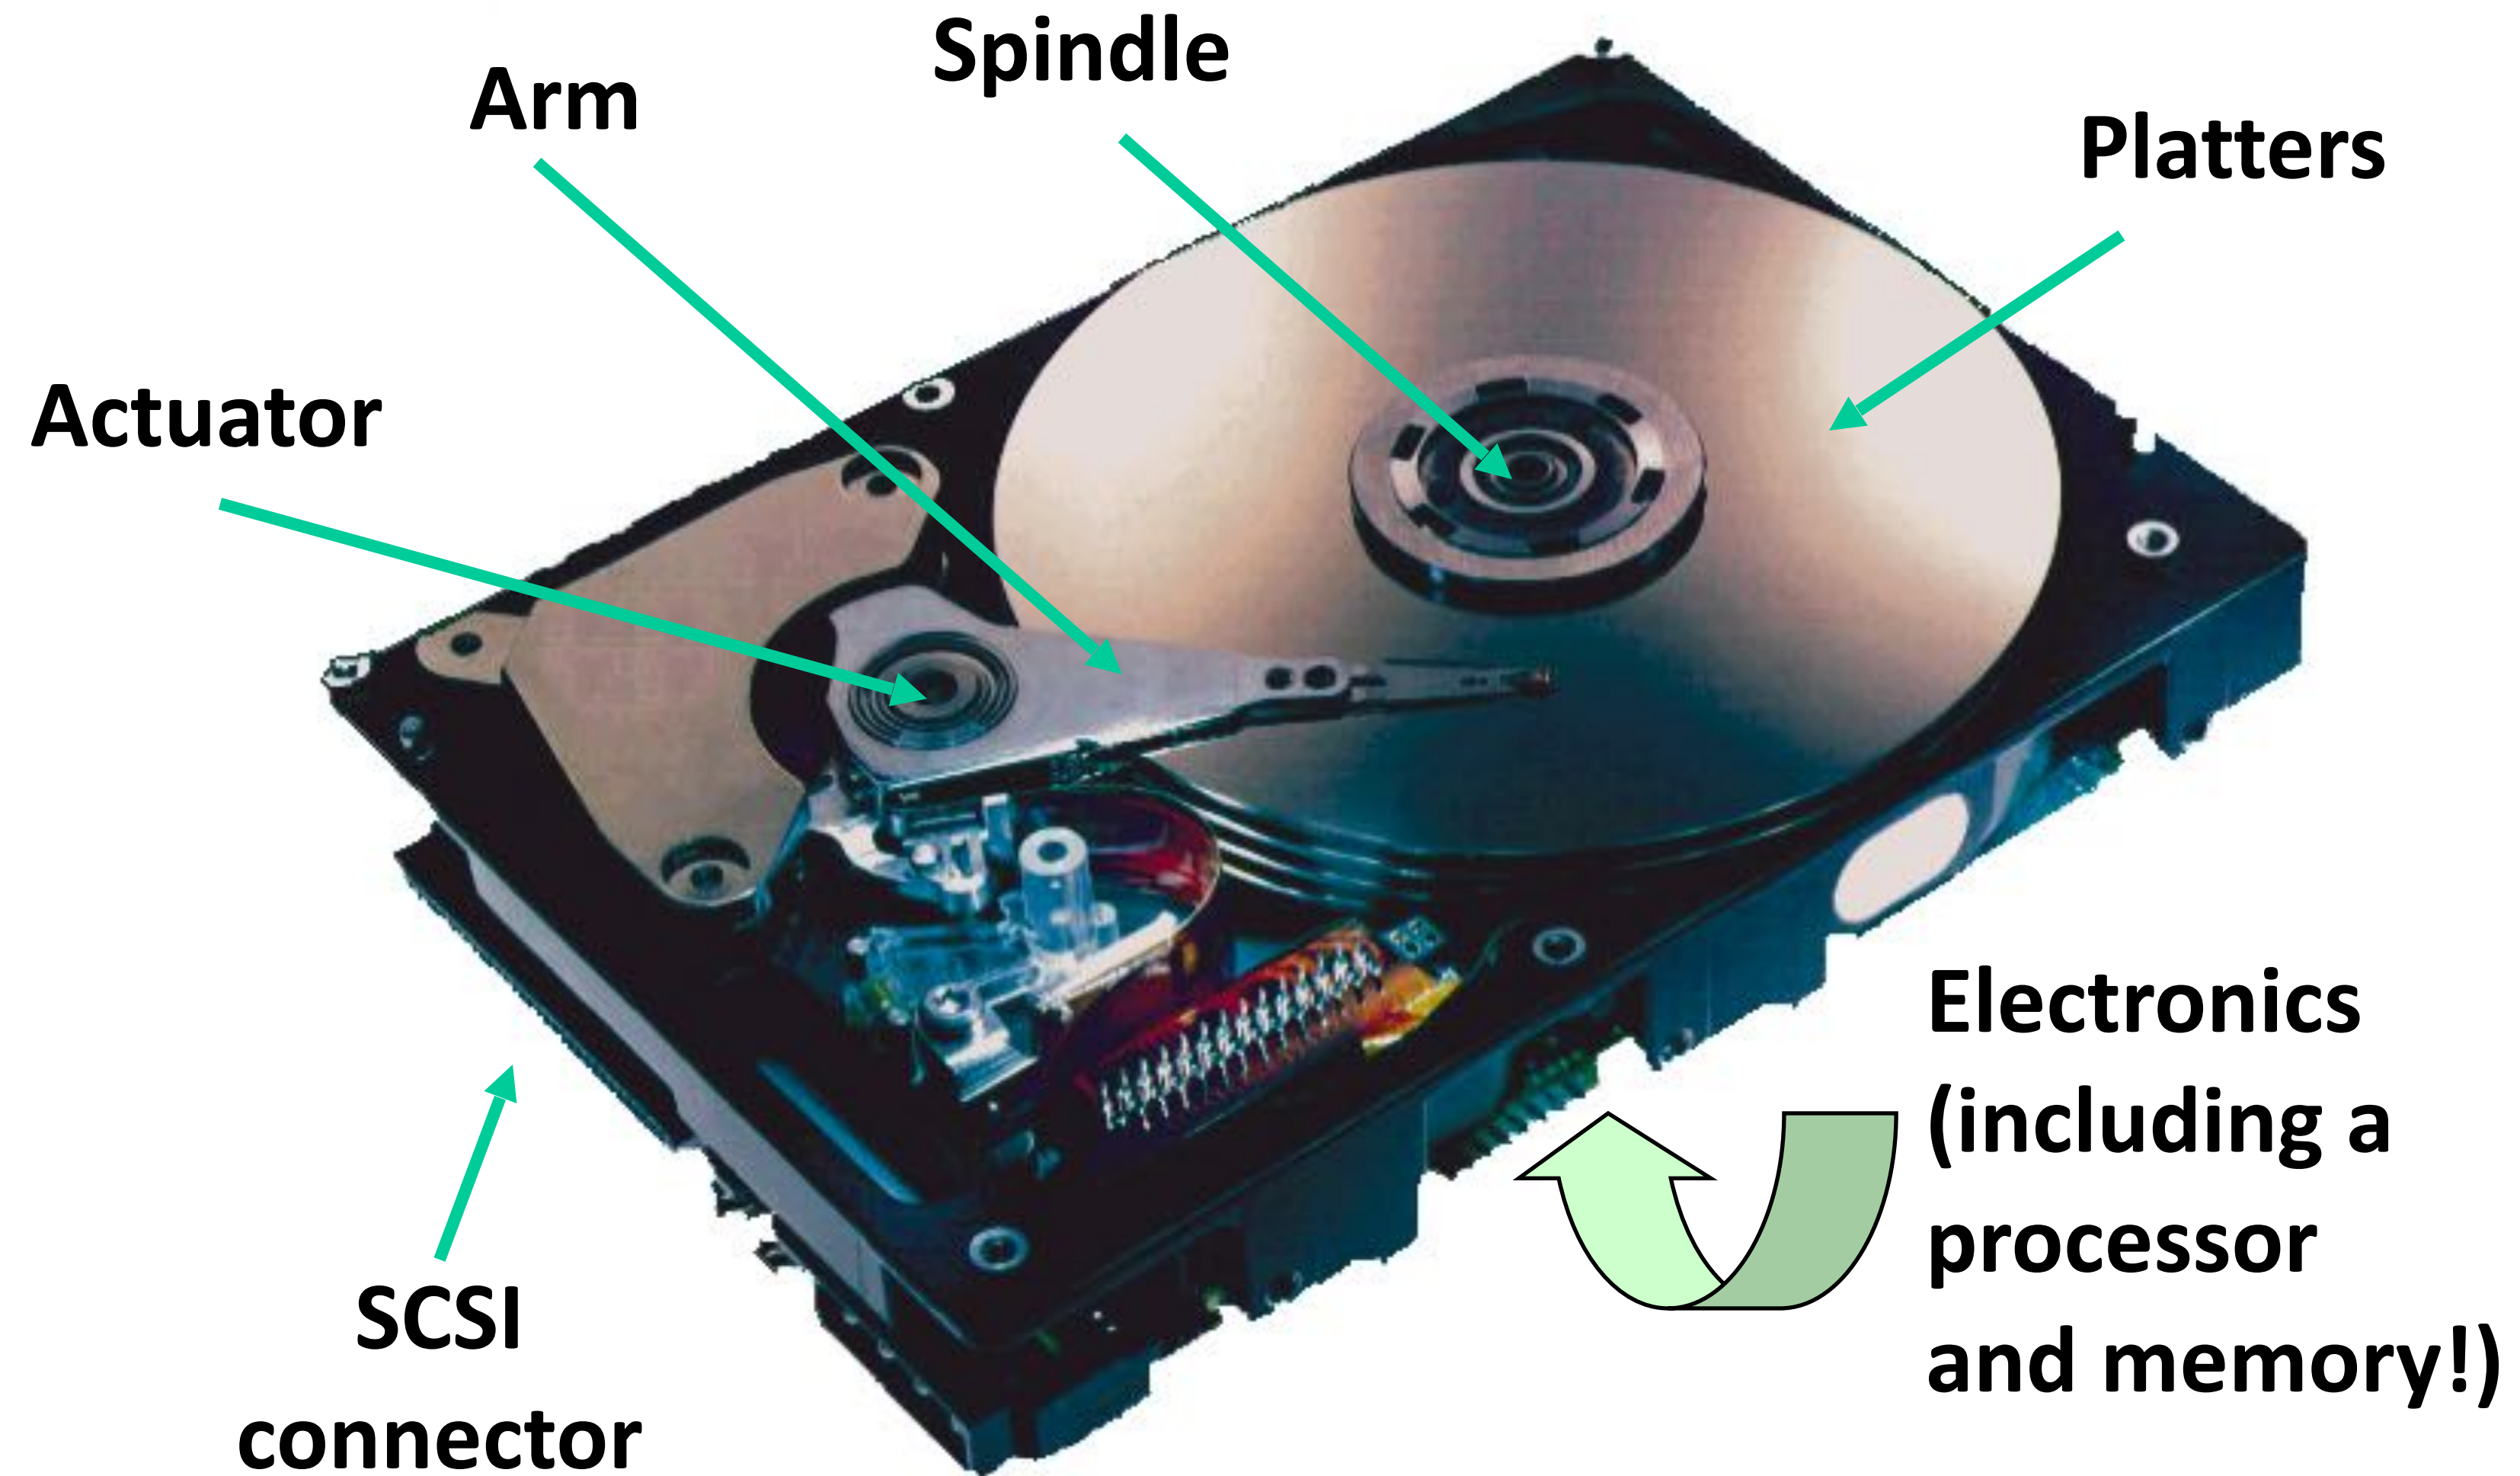



Image courtesy of Seagate Technology

Disk Geometry

- Disks consist of **platters**, each with two **surfaces**.
- Each surface consists of concentric rings called **tracks**.
- Each track consists of **sectors** separated by **gaps**.



Disk Capacity

Capacity: maximum number of bits that can be stored.


Vendors express capacity in units of gigabytes (GB) or terabytes (TB), where 1 GB = 10^9 Bytes and 1 TB = 10^{12} Bytes

Capacity is determined by these technology factors:

Recording density (bits/in): number of bits that can be squeezed into a 1 inch segment of a track.


Track density (tracks/in): number of tracks that can be squeezed into a 1 inch radial segment.

Areal density (bits/in²): product of recording and track density.



Disk Operation (Single-Platter View)

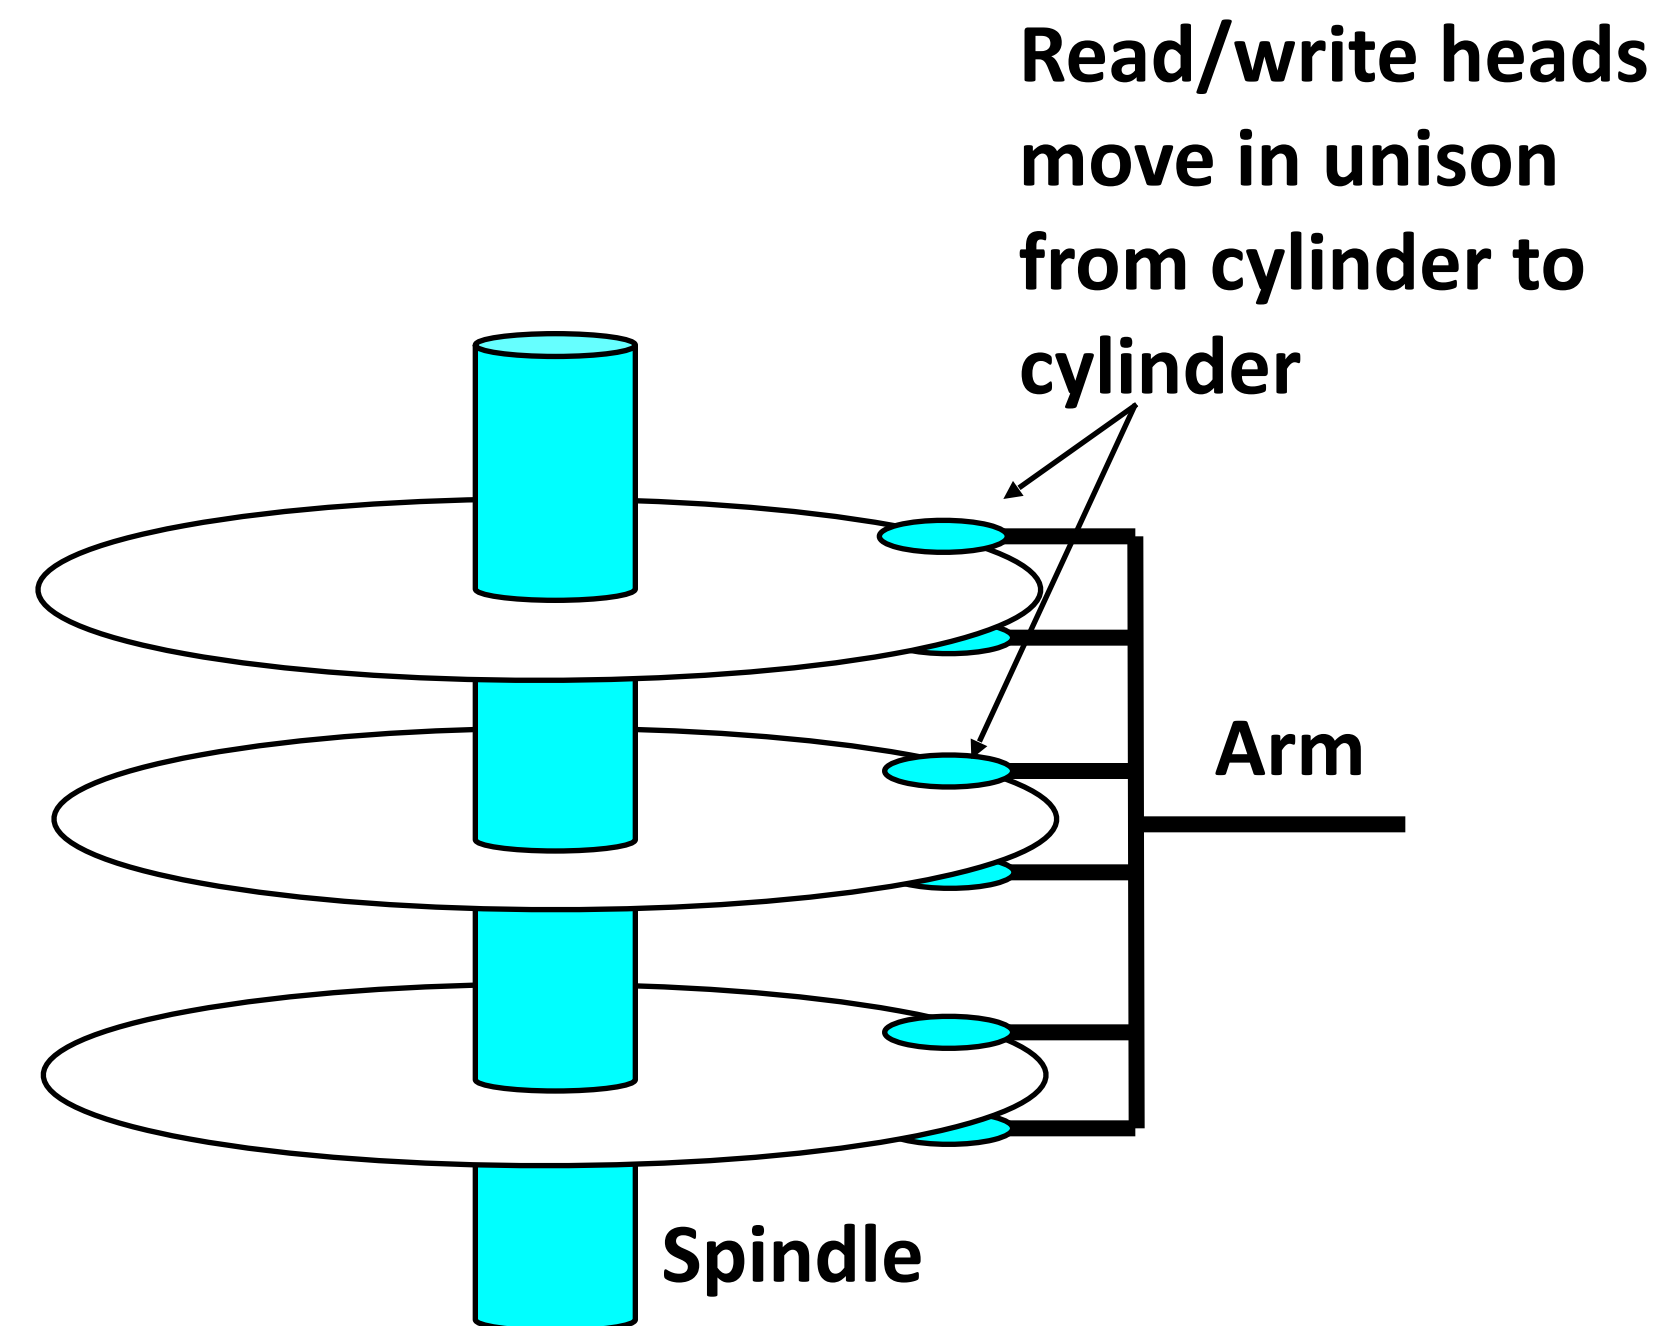
The disk surface spins at a fixed rotational rate




The read/write *head* is attached to the end of the *arm* and flies over the disk surface on a thin cushion of air.

By moving radially, the arm can position the read/write head over any track.

Disk Operation (Multi-Platter View)



Disk Access – Service Time Components



Disk Access Time

Average time to access some target sector approximated by:

$$T_{\text{access}} = T_{\text{avg seek}} + T_{\text{avg rotation}} + T_{\text{avg transfer}}$$

Seek time ($T_{\text{avg seek}}$)

Time to position heads over cylinder containing target sector.

Typical $T_{\text{avg seek}}$ is 3—9 ms

Rotational latency ($T_{\text{avg rotation}}$)

Time waiting for first bit of target sector to pass under r/w head.

$$T_{\text{avg rotation}} = 1/2 \times 1/\text{RPMs} \times 60 \text{ sec}/1 \text{ min}$$

Typical rotational rate = 7,200 RPMs

Transfer time ($T_{\text{avg transfer}}$)

Time to read the bits in the target sector.

$$T_{\text{avg transfer}} = 1/\text{RPM} \times 1/(\text{avg \# sectors/track}) \times 60 \text{ secs}/1 \text{ min}$$

time for one rotation (in minutes) fraction of a rotation to be read

Disk Access Time Example

Given:

Rotational rate = 7,200 RPM

Average seek time = 9 ms

Avg # sectors/track = 400

Derived:

$$T_{\text{avg rotation}} = 1/2 \times (60 \text{ secs}/7200 \text{ RPM}) \times 1000 \text{ ms/sec} = 4 \text{ ms}$$

$$T_{\text{avg transfer}} = 60/7200 \times 1/400 \times 1000 \text{ ms/sec} = 0.02 \text{ ms}$$

$$T_{\text{access}} = 9 \text{ ms} + 4 \text{ ms} + 0.02 \text{ ms}$$

Important points:

Access time dominated by seek time and rotational latency.

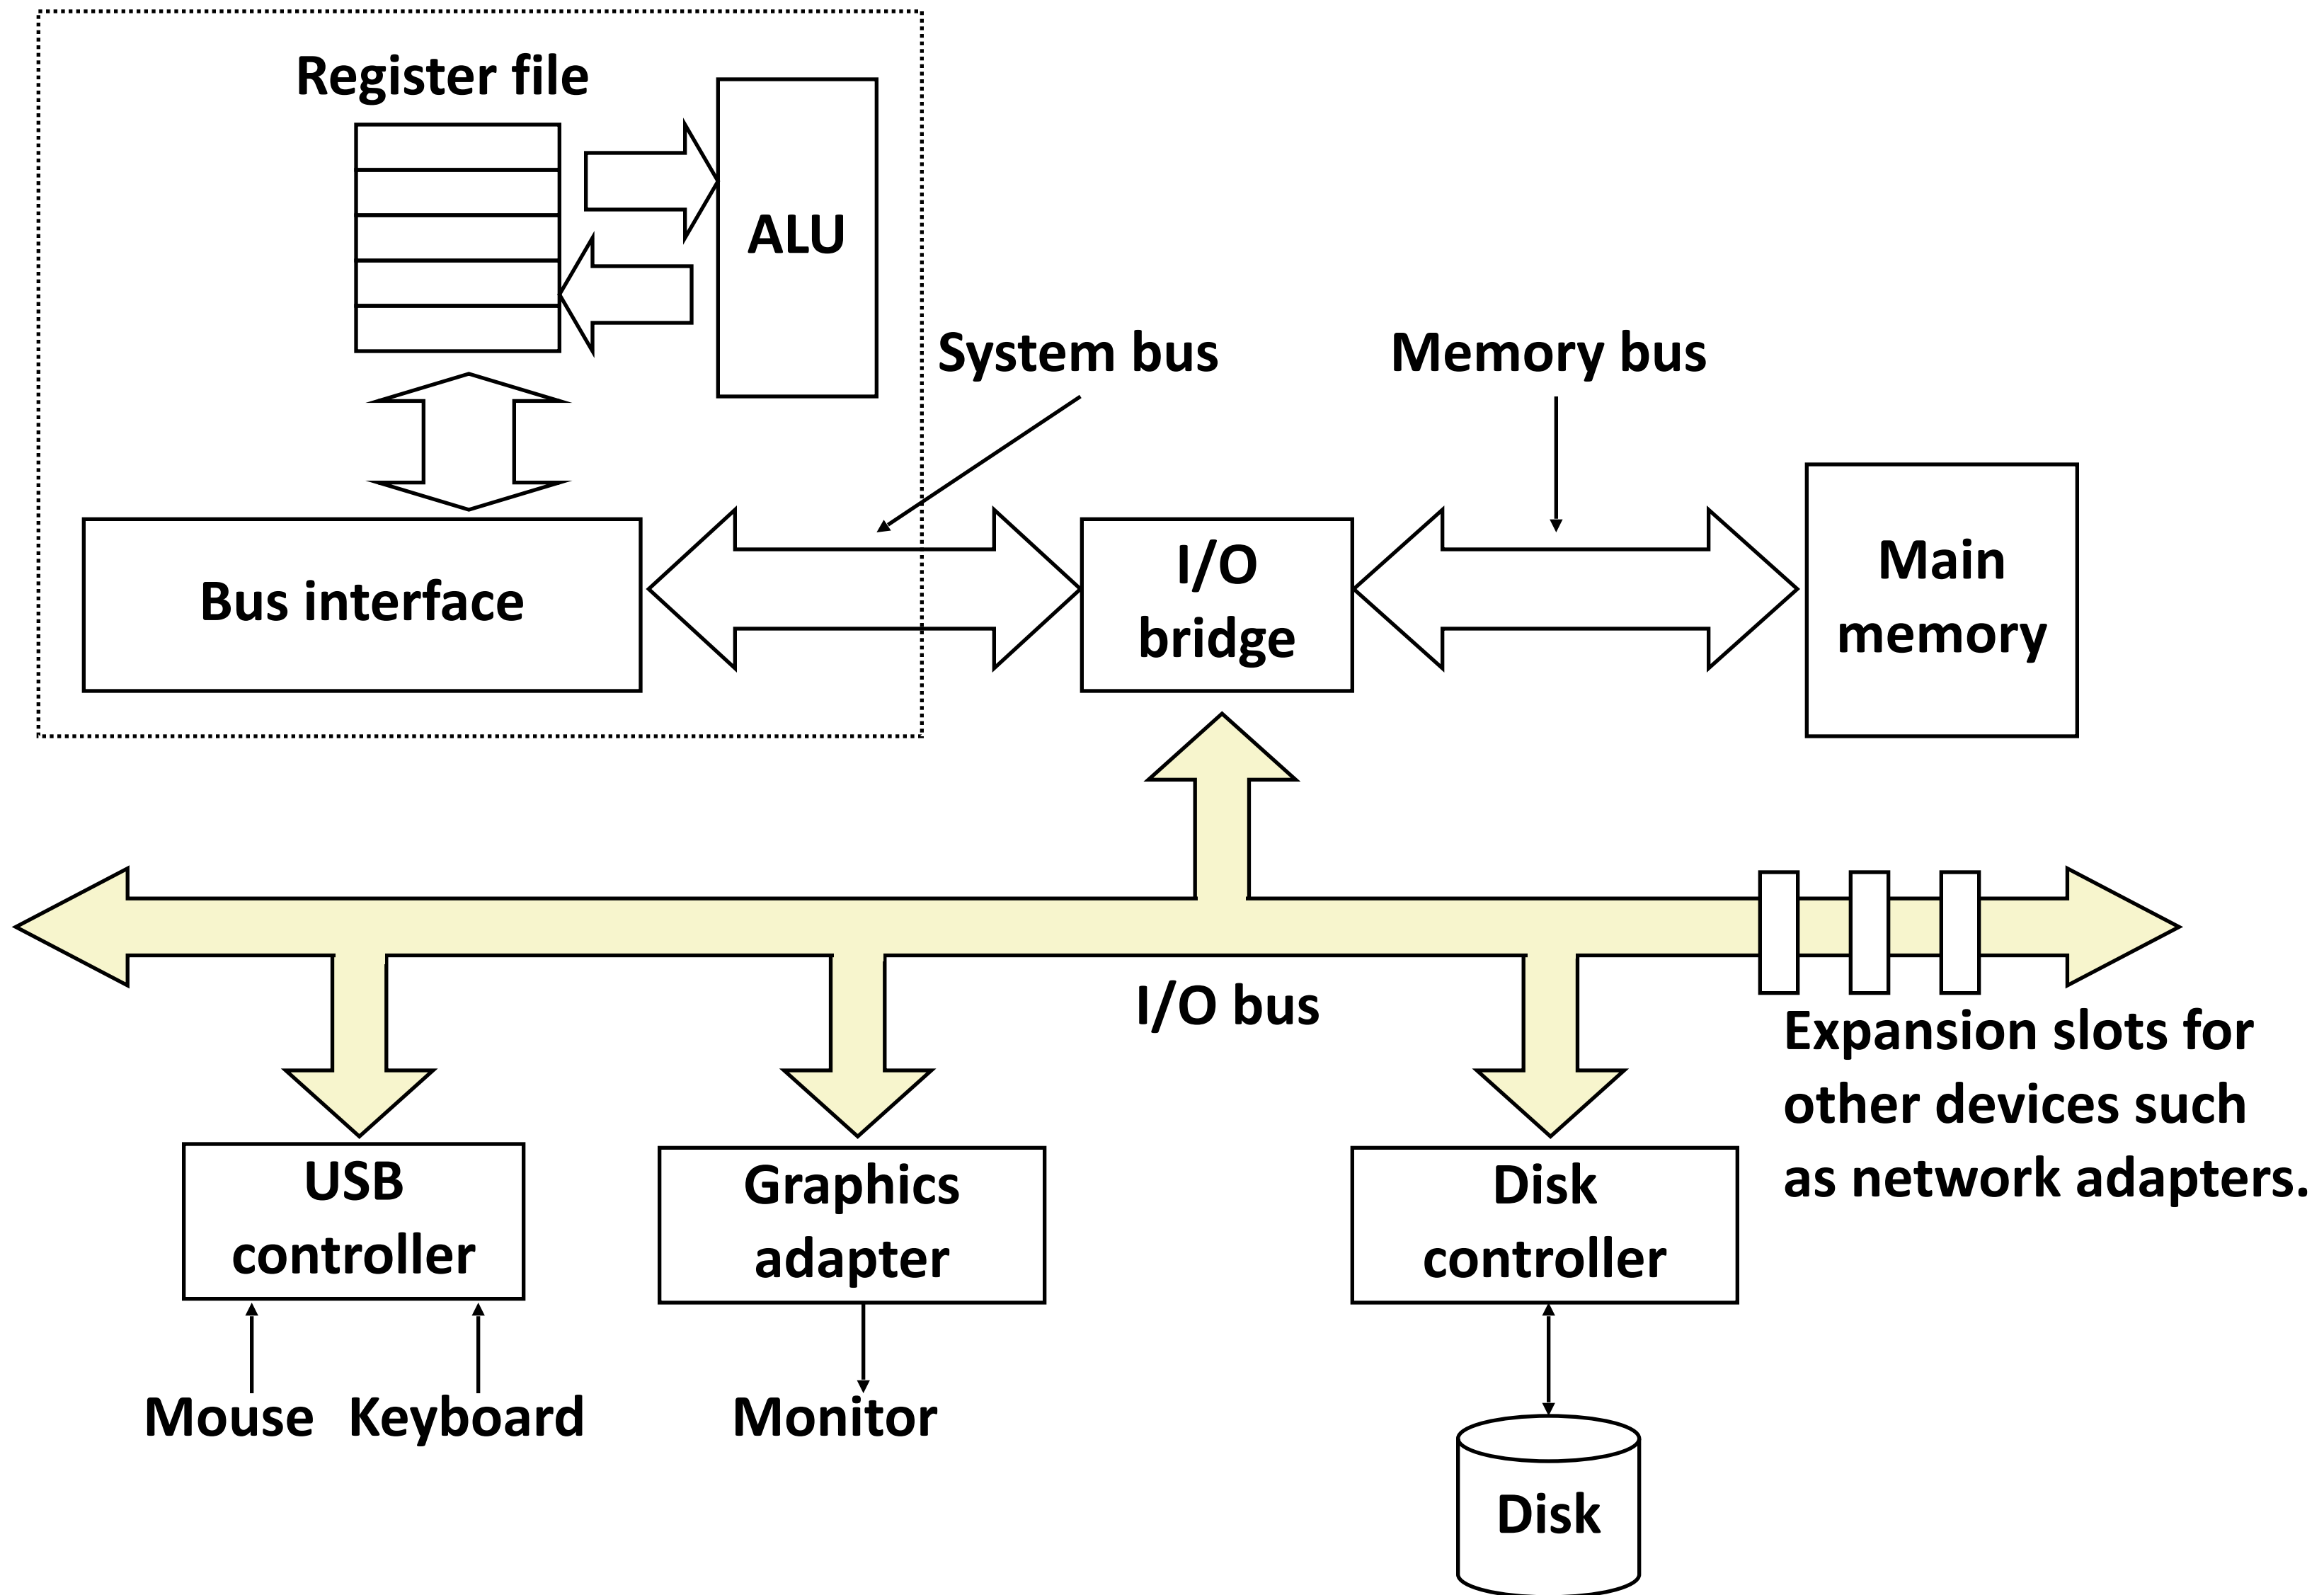
First bit in a sector is the most expensive, the rest are free.

SRAM access time is about 4 ns/doubleword, DRAM about 60 ns


Disk is about 40,000 times slower than SRAM,

2,500 times slower than DRAM.

I/O Bus




Reading a Disk Sector (1)




CPU initiates a disk read by writing a command, logical block number, and destination memory address to a **port** (address) associated with disk controller.

Reading a Disk Sector (2)



Disk controller reads the sector and performs a direct memory access (**DMA**) transfer into main memory.

Reading a Disk Sector (3)



When the DMA transfer completes, the disk controller notifies the CPU with an *interrupt* (i.e., asserts a special “interrupt” pin on the CPU).

Processor Performance

Q: *How fast can a processor process a program?*

- Modern CPUs can run millions of instructions per second!
 - ISA tells #**clock cycles** per instruction
 - CPU's **clock rate** helps map that to runtime (ns)
- Alas, most programs do not keep CPU always busy:
 - Memory access commands **stall** the processor; ALU and CU are *idle* during DRAM-register transfer
 - Worse, data may not be in DRAM—wait for (disk) I/O!
 - So, actual *runtime* of program may be OOM higher than what clock rate calculation suggests

Key Principle: Optimizing access to DRAM and use of processor caches is critical for processor performance!

Summary

- The speed gap between CPU, memory and mass storage continues to widen.
- Well-written programs exhibit a property called *locality*.
- Memory hierarchies based on *caching* close the gap by exploiting locality.

Review Questions

- What is an ISA?
- What are the 3 main kinds of commands in an ISA?
- Why do CPUs have both registers and caches?
- Why is it typically impossible for data processing programs to achieve 100% processor utilization?